

Prague AWS User Group

Databases on AWS: Purpose-built DBs for all your applications needs

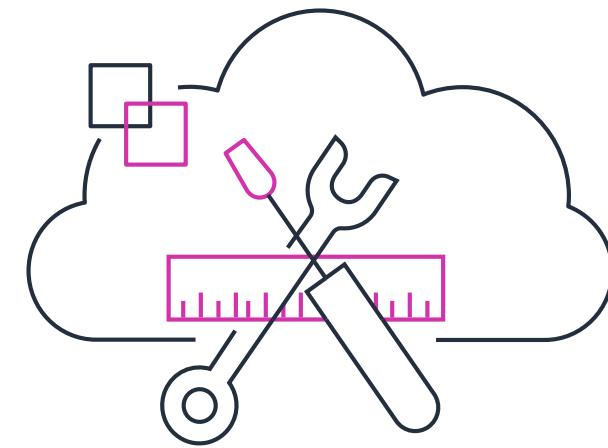
Vladimir Simek, Enterprise Solutions Architect
Prague AWS User Group
13th February 2019

How are you using your databases?

Two fundamental options



“Lift and shift” existing
apps to the cloud



Quickly build new
apps in the cloud

“Lift and shift” existing apps to the cloud



“Lift and shift” existing
apps to the cloud



Quickly build new
apps in the cloud

Amazon Relational Database Service (RDS)



Managed relational database service with a choice of six popular database engines

Amazon Aurora

MySQL®

PostgreSQL

MariaDB

Microsoft SQL Server®

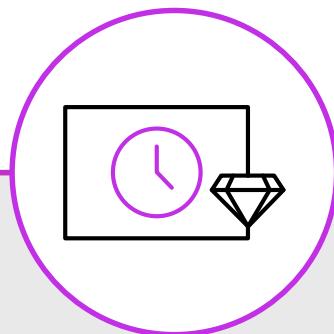
ORACLE®

Easy to administer



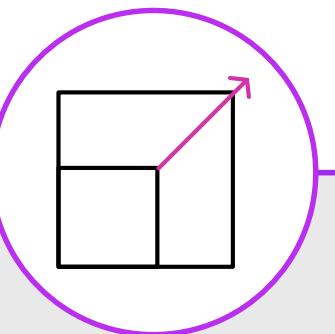
No need for infrastructure provisioning, installing, and maintaining DB software

Available and durable



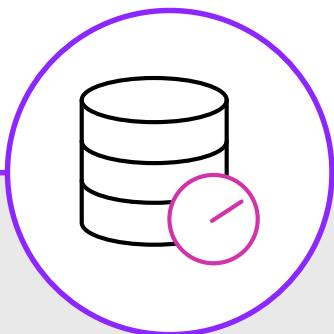
Automatic Multi-AZ data replication; automated backup, snapshots, failover

Highly scalable



Scale database compute and storage with a few clicks with no app downtime

Fast and secure



SSD storage and guaranteed provisioned I/O; data encryption at rest and in transit

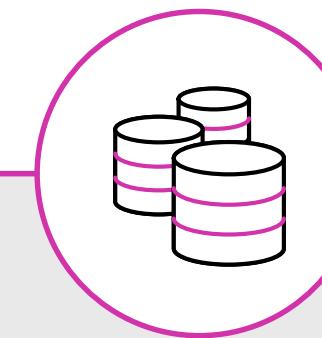
Amazon Aurora



MySQL and PostgreSQL-compatible relational database built for the cloud

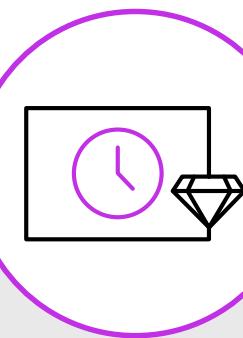
Performance and availability of commercial-grade databases at 1/10th the cost

Performance and scalability



5x throughput of standard MySQL and 3x of standard PostgreSQL; scale-out up to 15 read replicas

Availability and durability



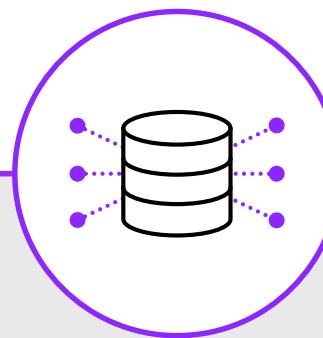
Fault-tolerant, self-healing storage; six copies of data across three Availability Zones; continuous backup to Amazon S3

Highly secure



Network isolation, encryption at rest/transit

Fully managed



Managed by RDS: No hardware provisioning, software patching, setup, configuration, or backups

AWS Database Migration Service (AWS DMS)



MIGRATING DATABASES TO AWS

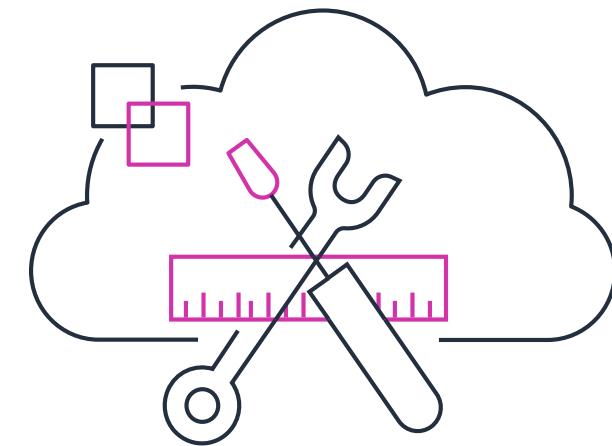
110,000+
databases migrated

-  Migrate between on-premises and AWS
-  Migrate between databases
-  Automated schema conversion
-  Data replication for zero-downtime migration

Quickly build new apps in the cloud



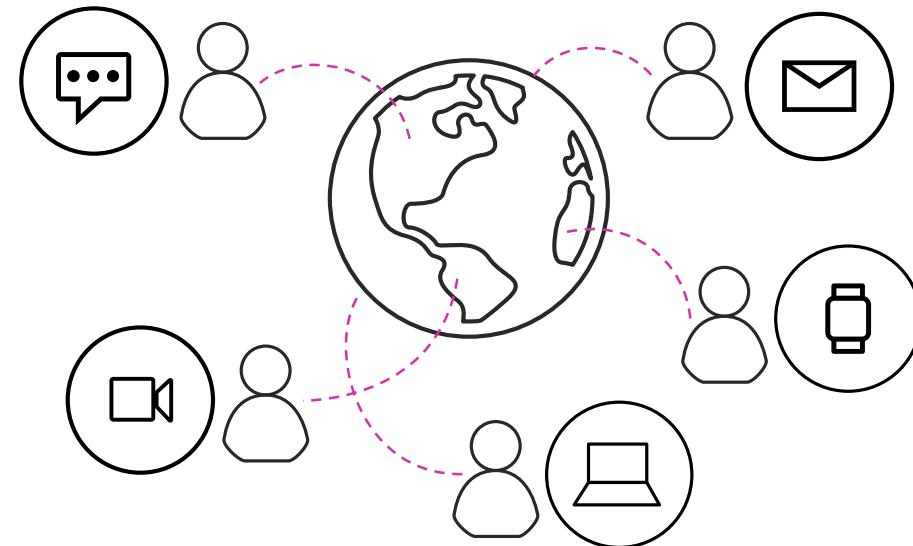
“Lift and shift” existing
apps to the cloud



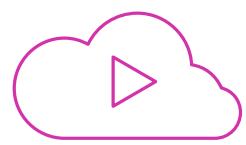
Quickly build new
apps in the cloud



Modern apps create new requirements



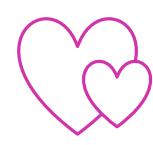
Ride hailing



Media streaming



Social media



Dating

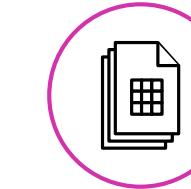
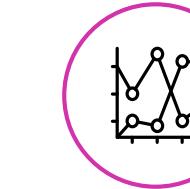
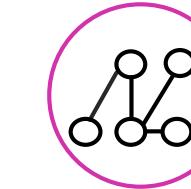
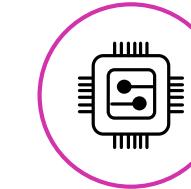
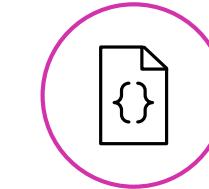
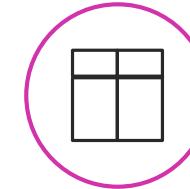
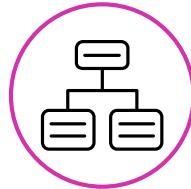
- Users: 1 million+
- Data volume: TB–PB–EB
- Locality: Global
- Performance: Milliseconds–microseconds
- Request rate: Millions
- Access: Web, mobile, IoT, devices
- Scale: Up-down, Out-in
- Economics: Pay for what you use
- Developer access: No assembly required

One size doesn't fit all





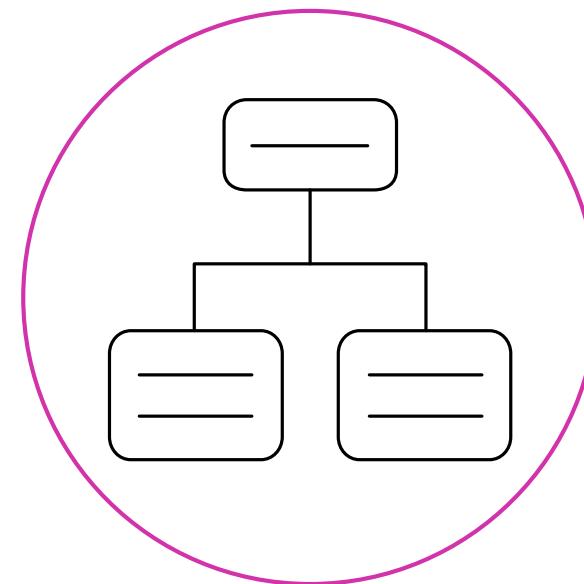
Common data categories and use cases



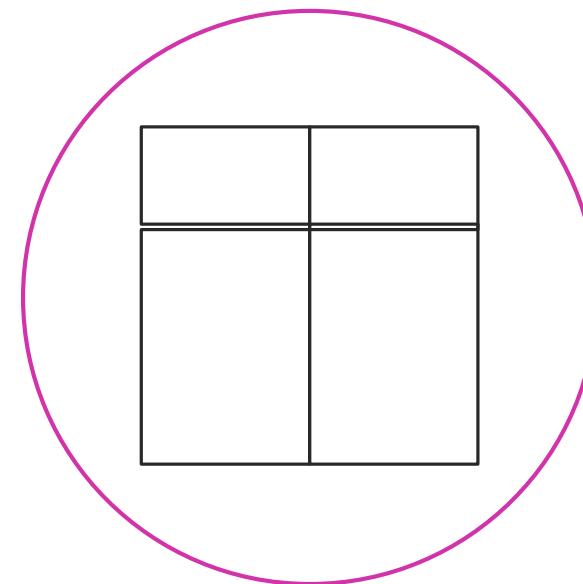
Relational	Key-value	Document	In-memory	Graph	Time-series	Ledger
Referential integrity, ACID transactions, schema-on-write	High throughput, low-latency reads and writes, endless scale	Store documents and quickly access querying on any attribute	Query by key with microsecond latency	Quickly and easily create and navigate relationships between data	Collect, store, and process data sequenced by time	Complete, immutable, and verifiable history of all changes to application data
Lift and shift, ERP, CRM, finance	Real-time bidding, shopping cart, social, product catalog, customer preferences	Content management, personalization, mobile	Leaderboards, real-time analytics, caching	Fraud detection, social networking, recommendation engine	IoT applications, event tracking	Systems of record, supply chain, health care, registrations, financial



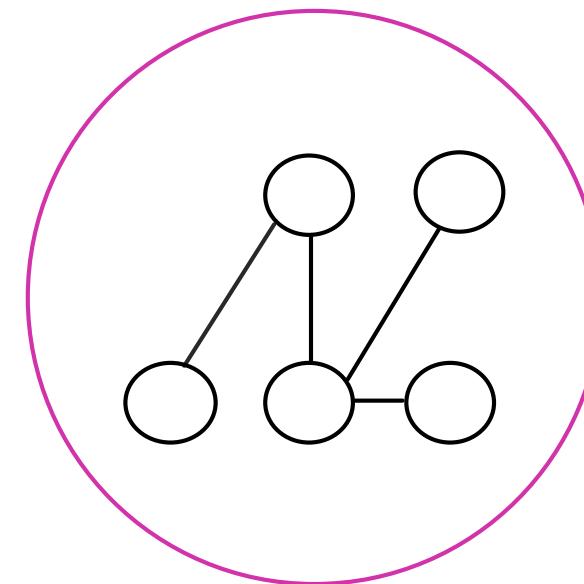
Let's take a closer look at...



Relational



Key-value

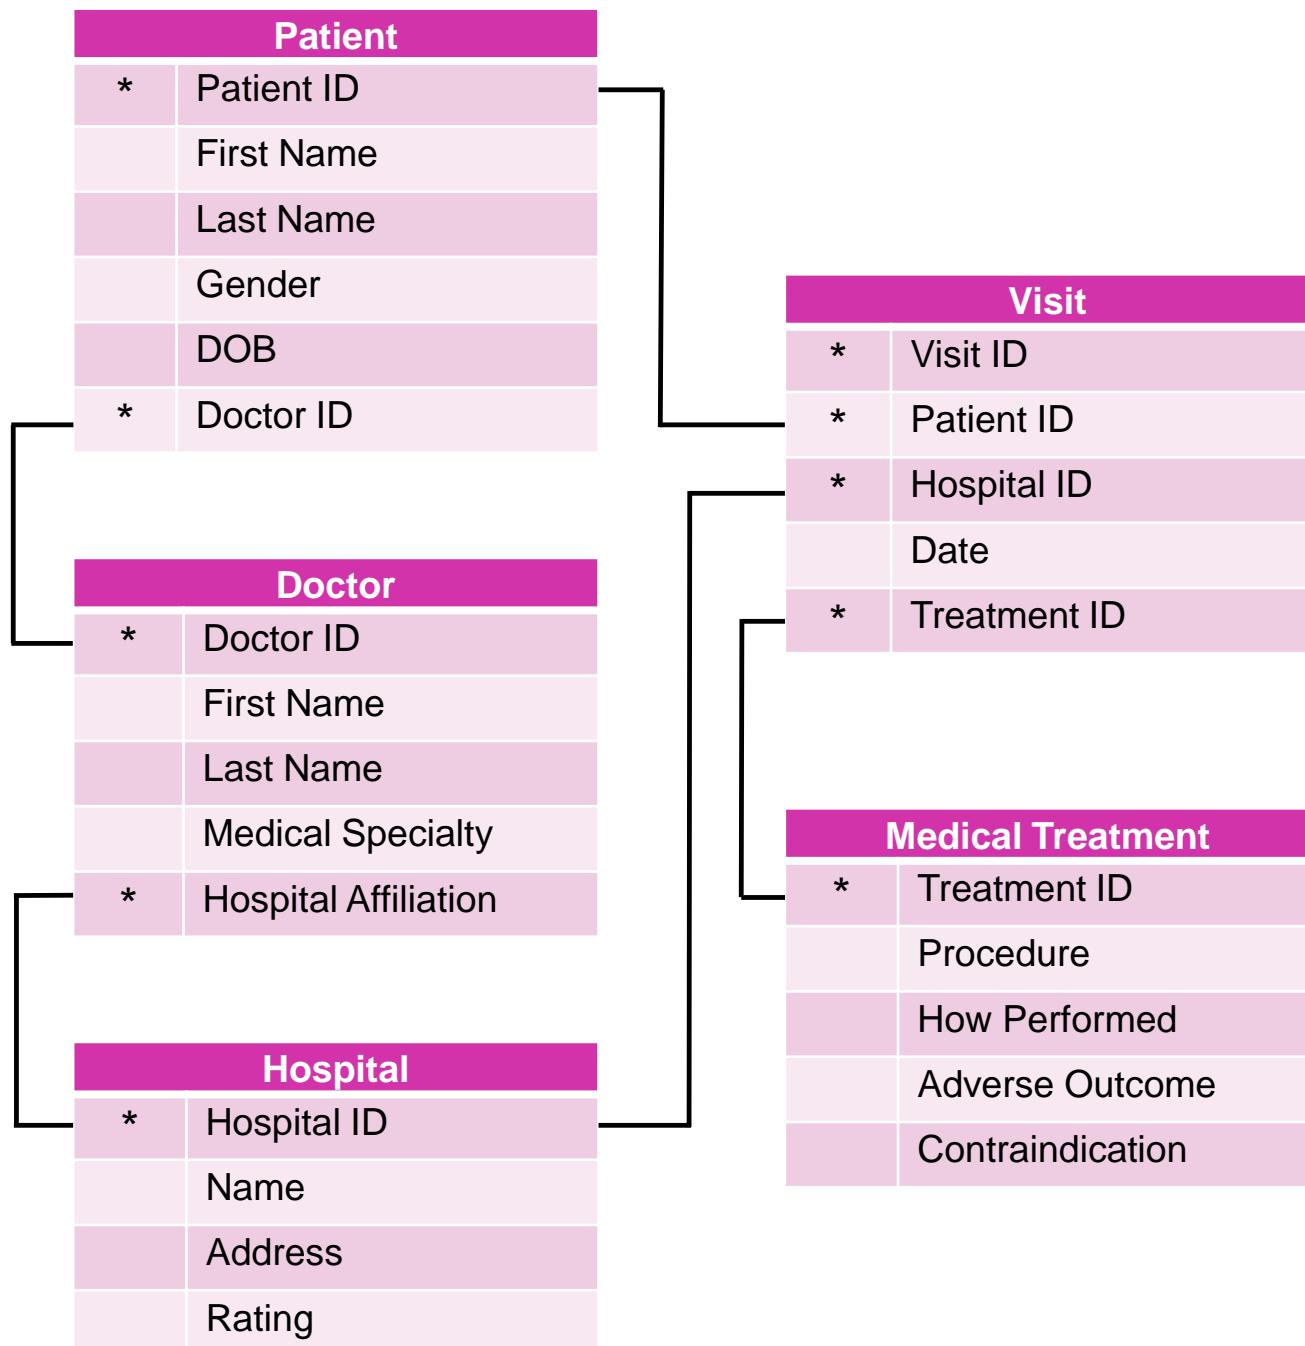


Graph



Relational data

- Divide data among tables
- Highly structured
- Relationships established via keys enforced by the system
- Data accuracy and consistency

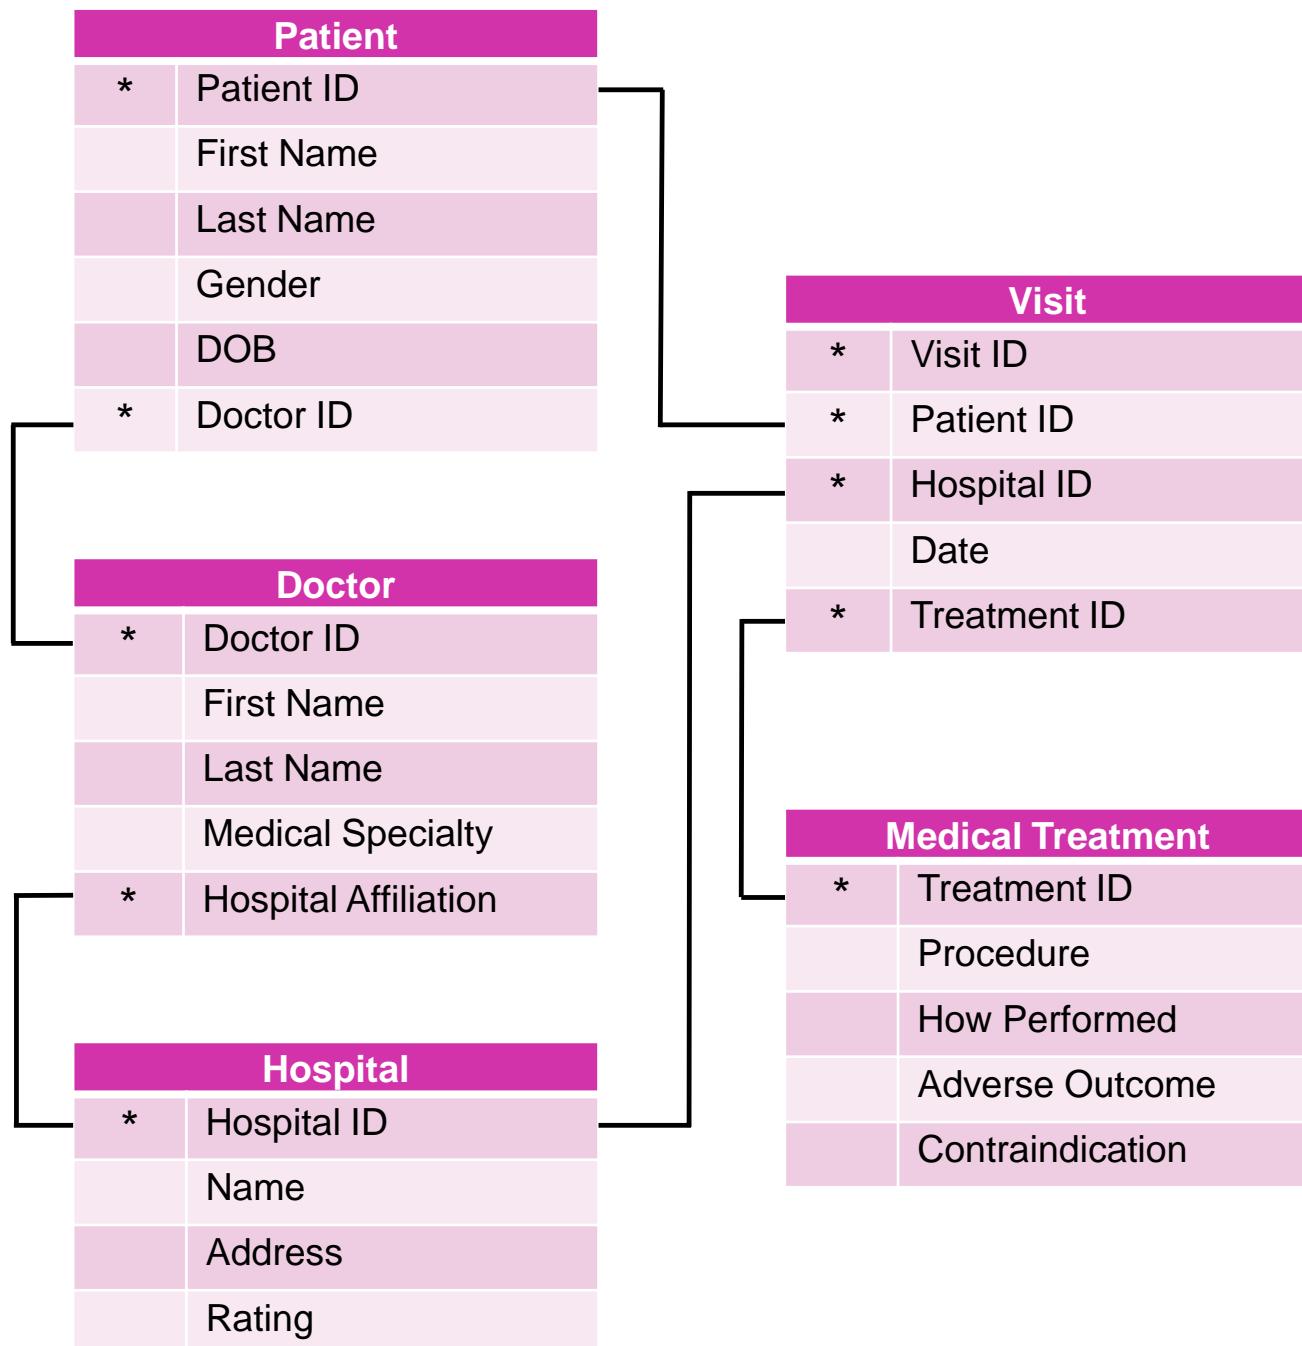




Relational use case

// Doctor affiliation with Mercy doctor hospitalized last week

```
SELECT
  SELECT d.first_name, d.last_name
  FROM doctor as d,
       visit as v,
       hospital as h,
       doctor as d
  WHERE d.hospital = h.hospital_id
    AND h.hospital_id = 'Mercy'
    AND h.hospital_id = d.hospital
    AND v.t_date > date_trunc('week',
      CURRENT_TIMESTAMP - interval '1 week')
  GROUP BY
    d.first_name, d.last_name;
```





Key-value data

- Simple key value pairs
- Partitioned by keys
- Resilient to failure
- High throughput, low-latency reads and writes
- Consistent performance at scale

```
PUT {
  TableName:"Gamers",
  Item: {
    "GamerTag": "Hammer57",
    "Level": 21,
    "Points": 4050,
    "Score": 483610,
    "Plays": 1722
  }
}
```

```
GET {
  TableName:"Gamers",
  Key: {
    "GamerTag": "Hammer57",
    "ProjectionExpression": "Points"
  }
}
```

Gamers				
Primary Key	Attributes			
GamerTag	Level	Points	High Score	Plays
Hammer57	21	4050	483610	1722
FluffyDuffy	5	1123	10863	43
Lol777313	14	3075	380500	1307
Jam22Jam	20	3986	478658	1694
ButterZZ_55	7	1530	12547	66
...



Key-value use case

// Status of Hammer57

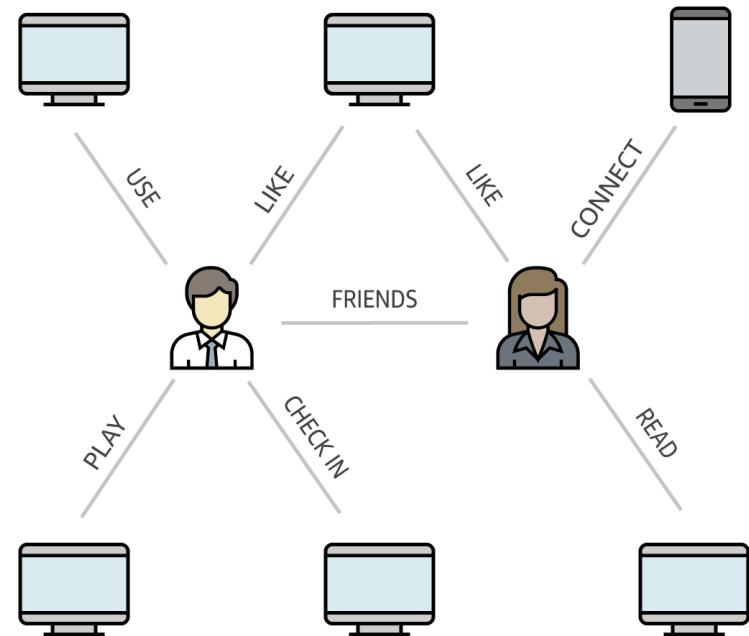
```
GET {  
  TableName:"Gamers",  
  Key: {  
    "GamerTag":"Hammer57",  
    "Type":"Status" } }
```

// Return all Hammer57

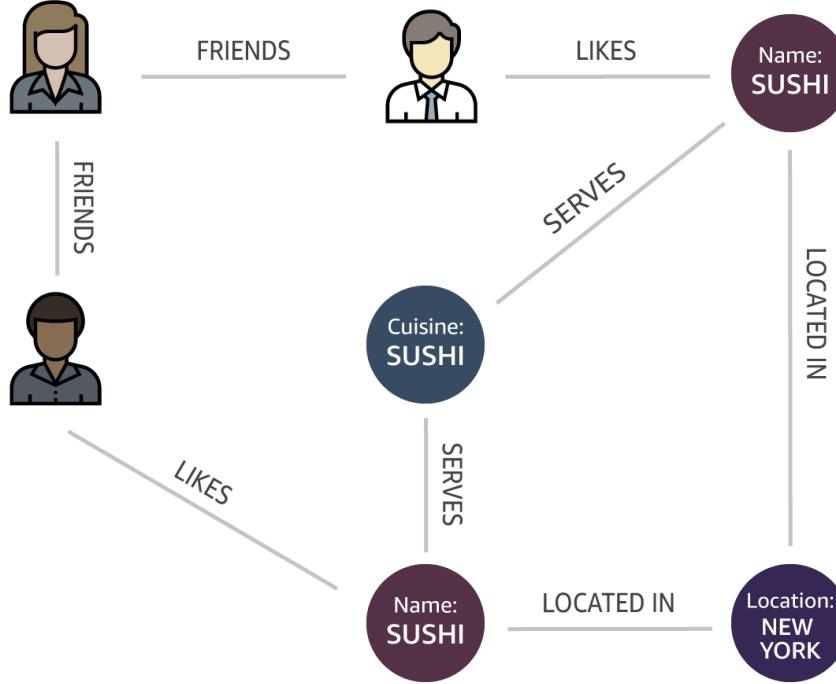
```
QUERY {  
  TableName:"Gamers",  
  KeyConditionExpression:"GamerTag = :a",  
  ExpressionAttributevalues: {  
    ":a":"Hammer57" } }
```

Gamers				
Primary Key		Attributes		
Gamer Tag	Type	Level	Points	Tier
Hammer57	Rank	87	4050	Elite
	Status	Health	Progress	
	Weapon	90	30	
FluffyDuffy	Rank	Class	Damage	Range
	Weapon	Taser	87%	50
	Rank	Level	Points	Tier
FluffyDuffy	Status	5	1072	Trainee
	Health	Progress		
	Weapon	37	8	

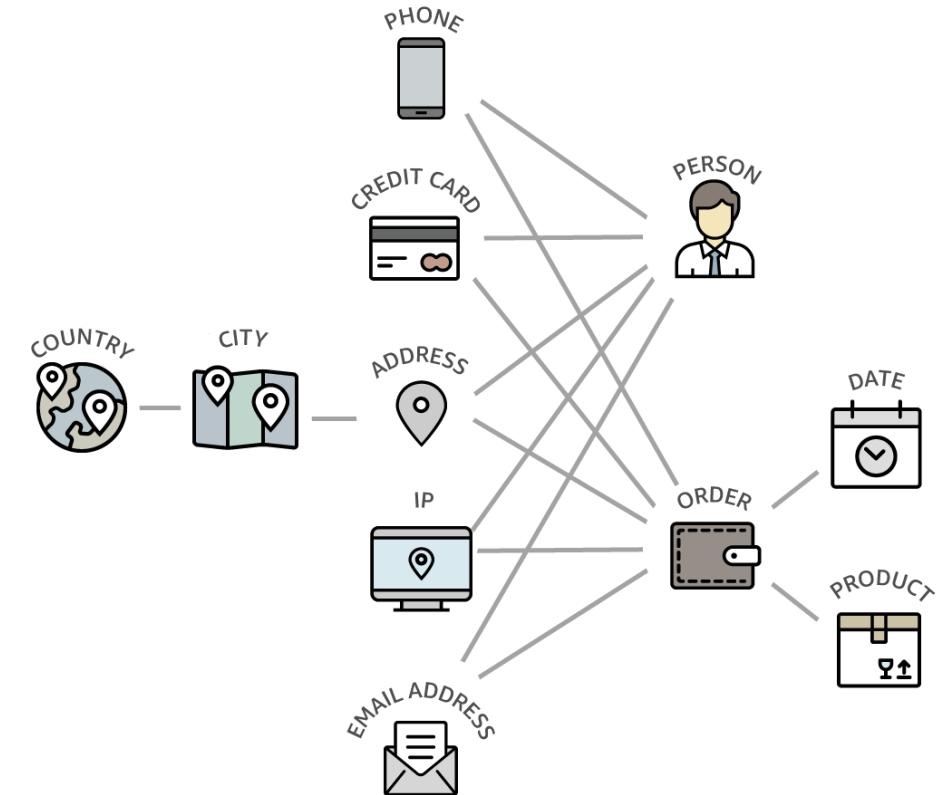
Relationships enable new applications



Social networks



Restaurant recommendations

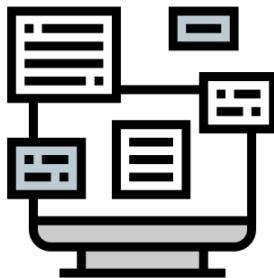


Retail fraud detection

Use cases for highly connected data



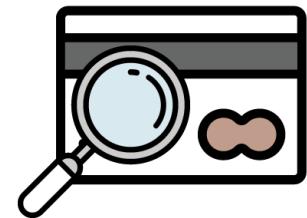
Social networking



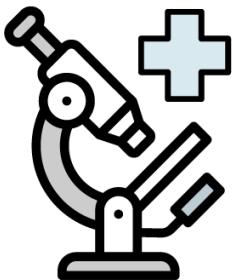
Recommendations



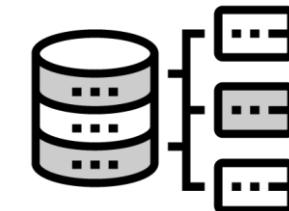
Knowledge graphs



Fraud detection



Life Sciences

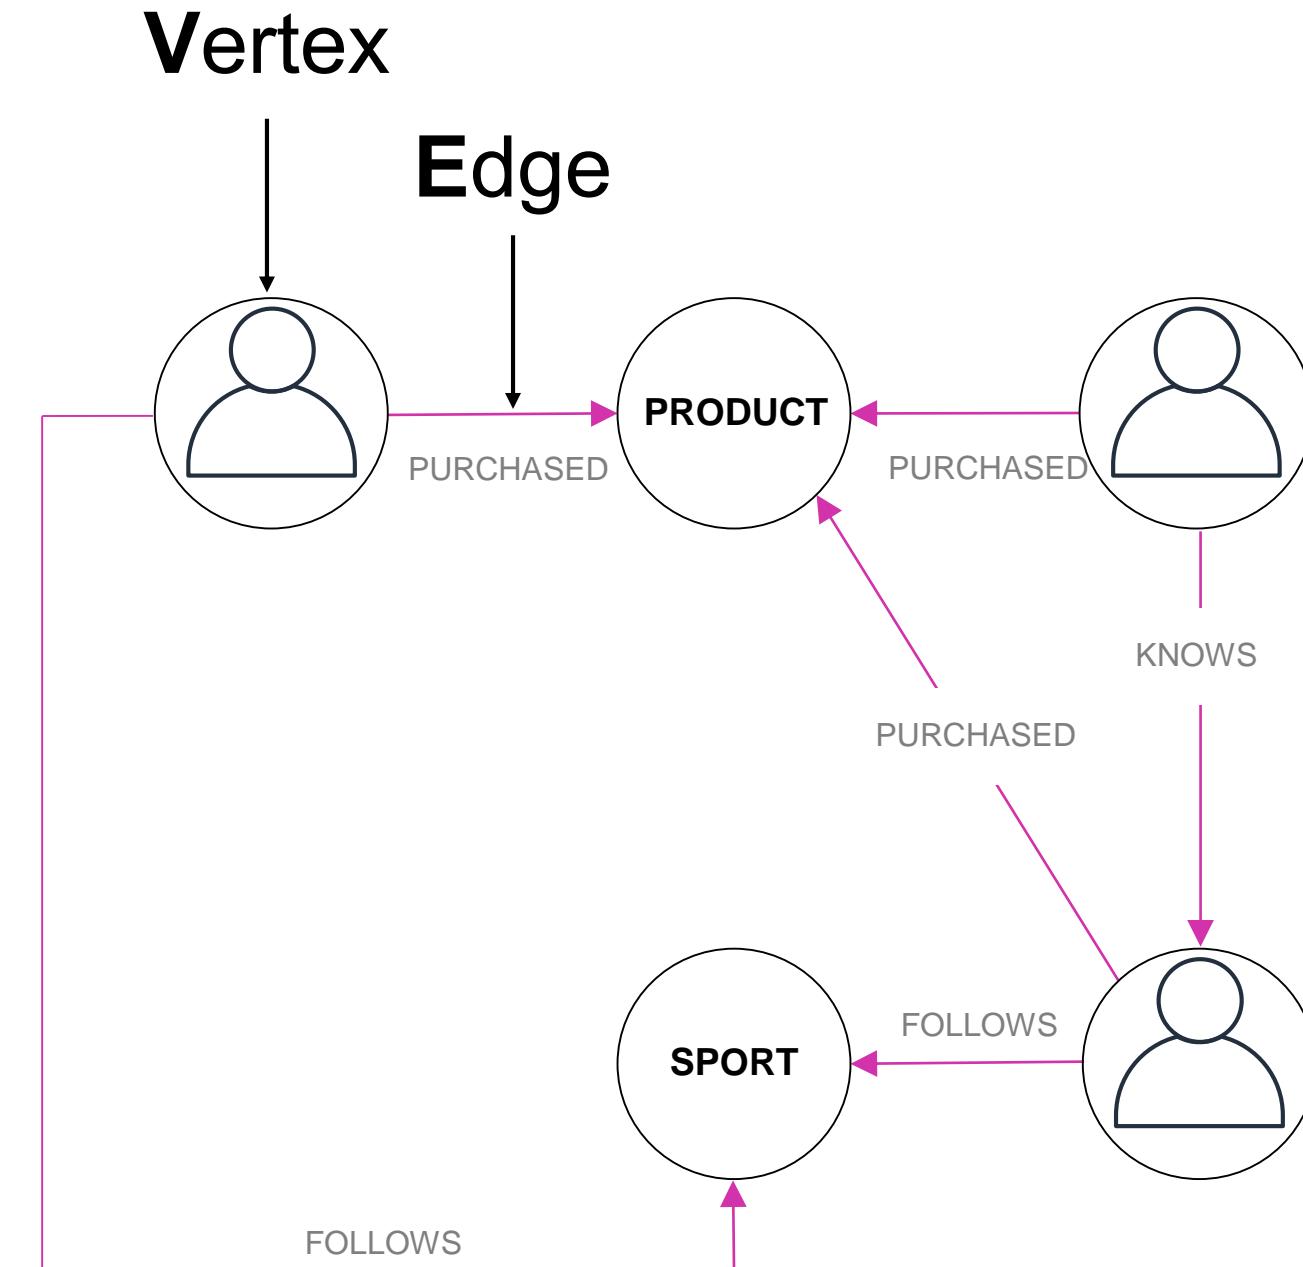


Network & IT operations



Graph data

- Relationships are first-class objects
- Vertices connected by Edges





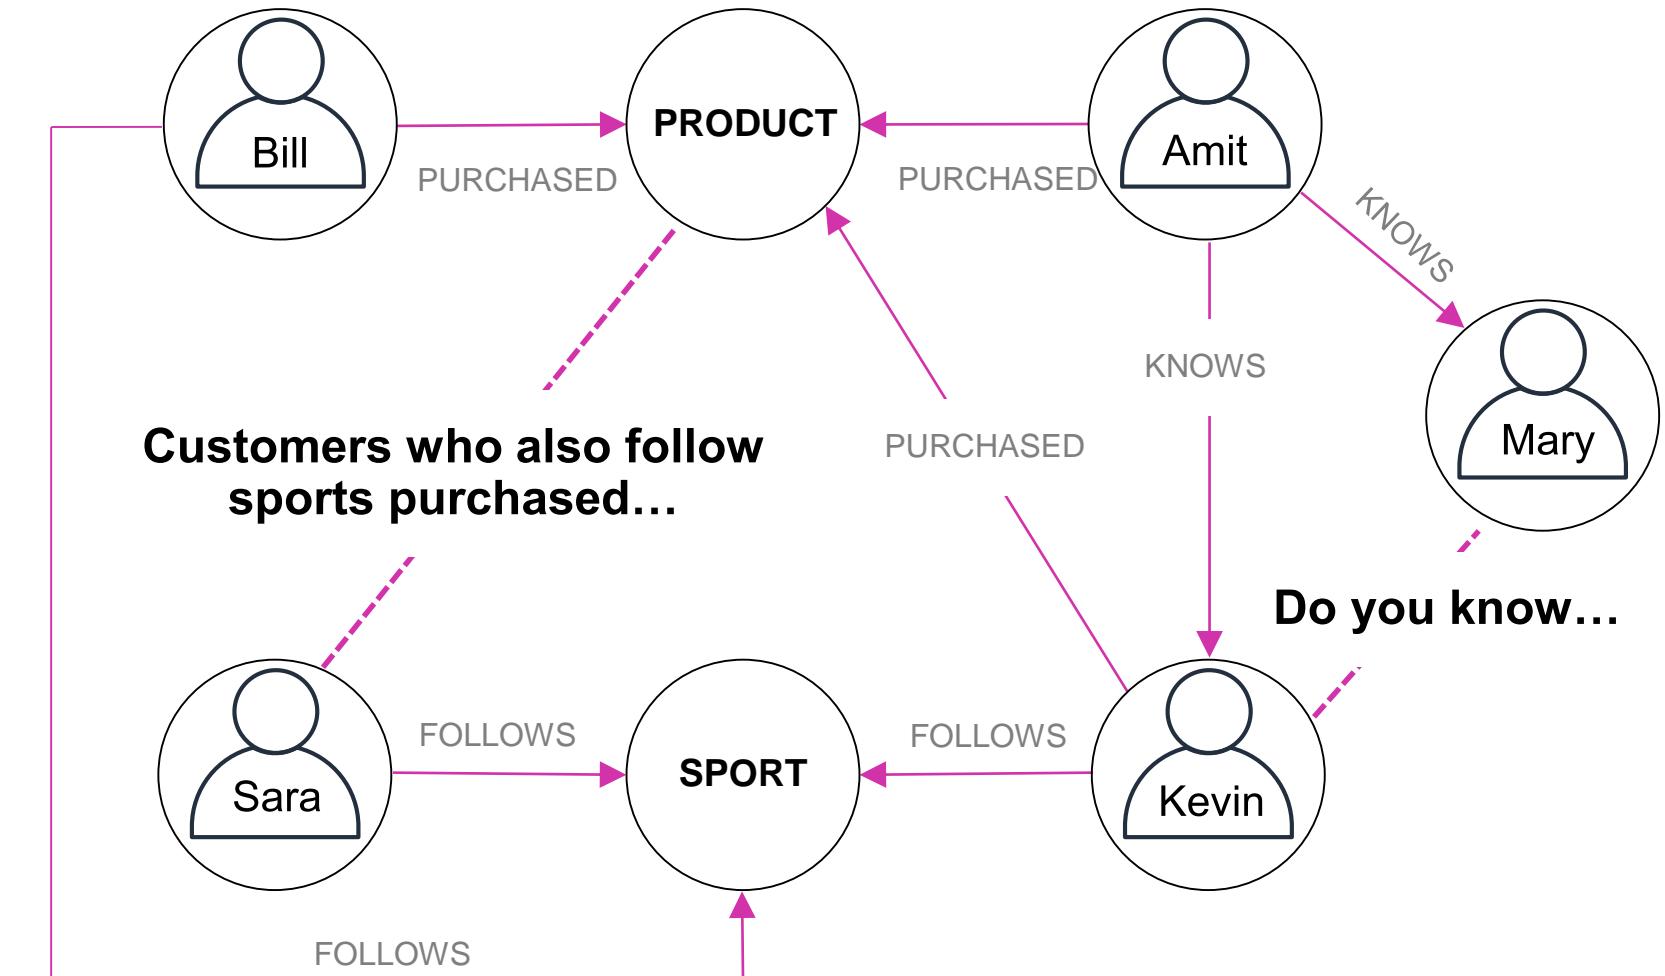
Graph use case

// Product recommendation to a user

```
gremlin> g.v().has('name', 'sara').as('customer').out('follows').in('follows').out('purchased')  
where(neq('customer')).dedup().by('name').properties('name')
```

// Identify a friend in common and make a recommendation

```
gremlin> g.v().has('name', 'mary').as('start').  
both('knows').both('knows').  
where(neq('start')).  
dedup().by('name').properties('name')
```





Airbnb uses different databases based on the purpose

User search history: **Amazon DynamoDB**

- Massive data volume
- Need quick lookups for personalized search

Session state: **Amazon ElastiCache**

- In-memory store for submillisecond site rendering

Relational data: **Amazon RDS**

- Referential integrity
- Primary transactional database



Spanish ▾



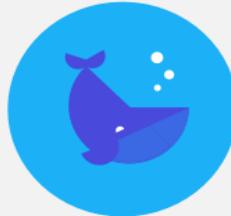
Basics



Basics 2



Phrases



Animals

0/6



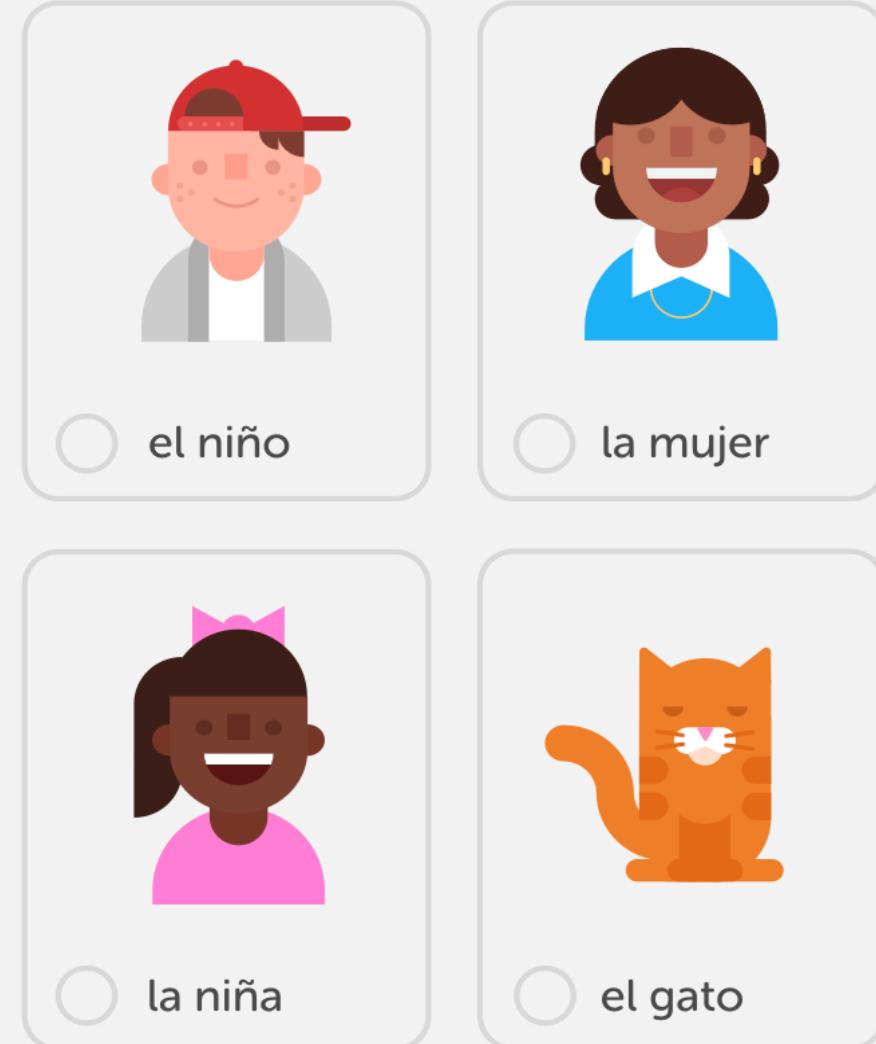
Clothing

0/6



Plurals

0/4



CHECK



300M total users

7B exercises per month

CHALLENGE

Wanted to enable anyone to learn a language for free.

SOLUTION

Purpose-built databases from AWS:

- **DynamoDB:** 31B items tracking which language exercises completed
- **Aurora:** primary transactional database for user data
- **ElastiCache:** instant access to common words and phrases

Result:

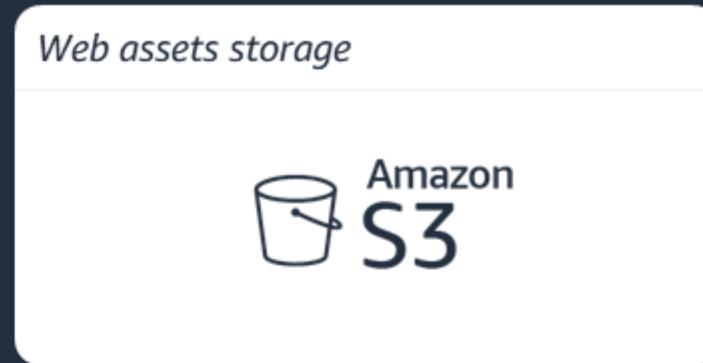
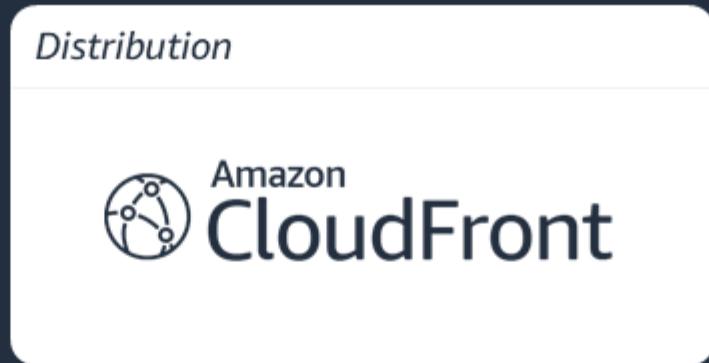
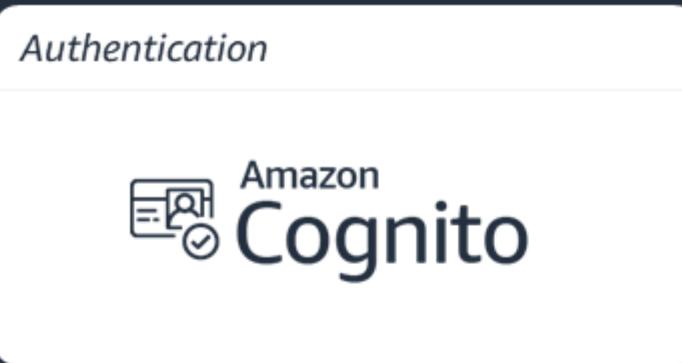
More people learning a language on Duolingo than entire US school system



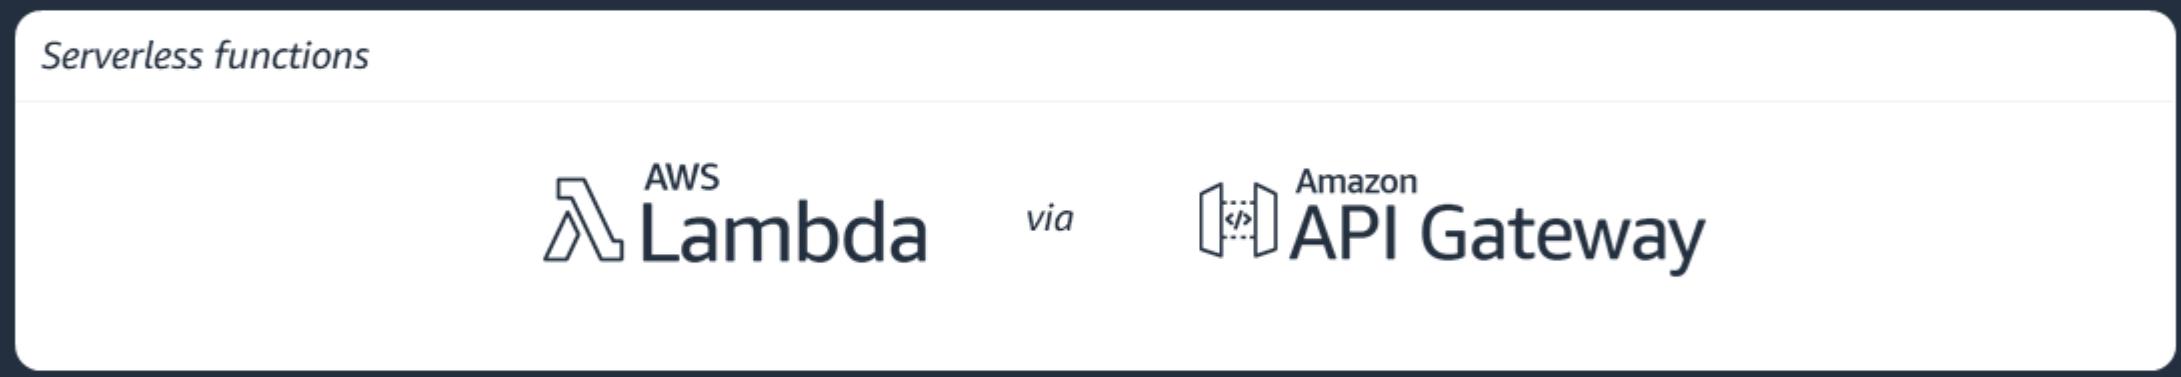
Demo

Demo App Architecture

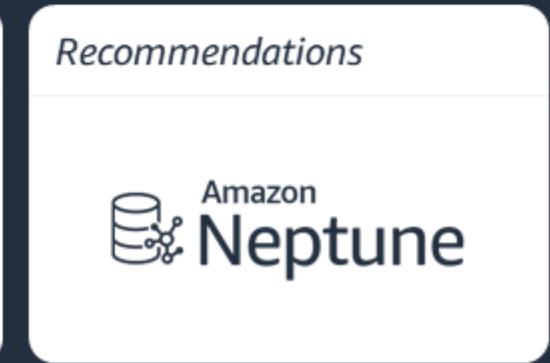
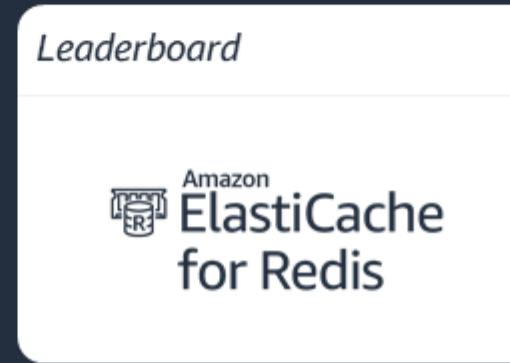
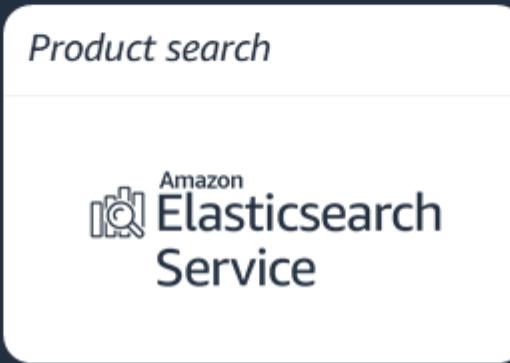
Frontend



Interface

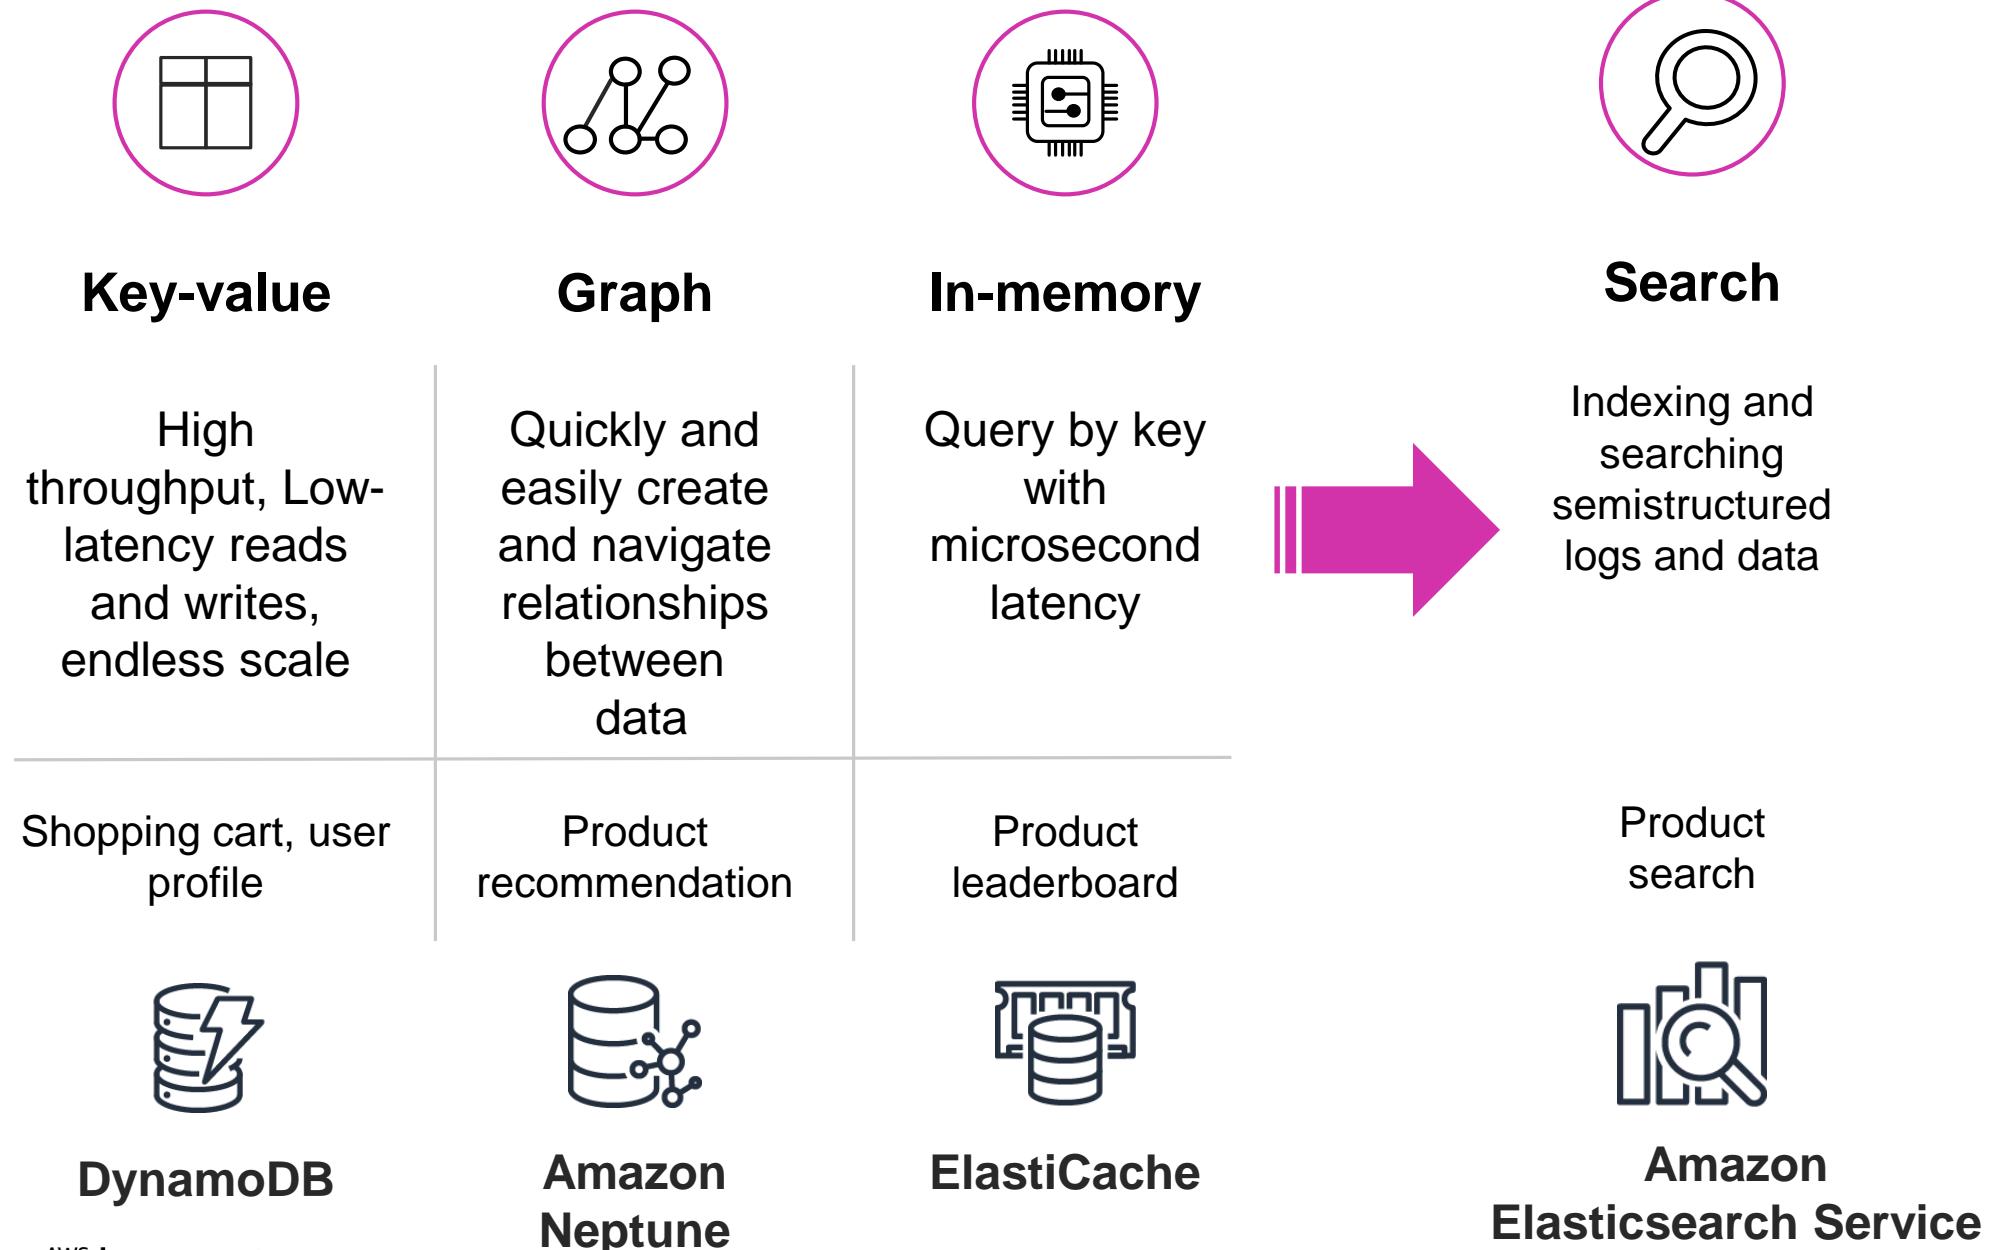


Backend





Retail demo application



Demo application:

1. Available today

2. On GitHub:

[/aws-samples/aws-bookstore-demo-app](https://github.com/aws-samples/aws-bookstore-demo-app)

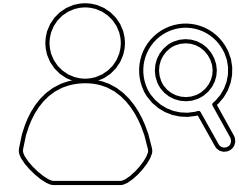
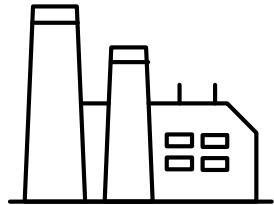
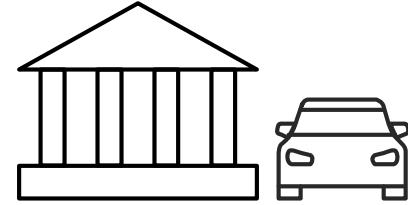
3. One click CloudFormation deployment

Ledger database



Common customer use cases

Ledgers with centralized control



Healthcare

Verify and track hospital equipment inventory

Government

Track vehicle title history

Manufacturers

Track distribution of a recalled product

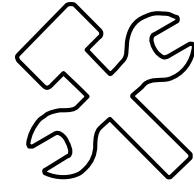
HR & Payroll

Track changes to an individual's profile



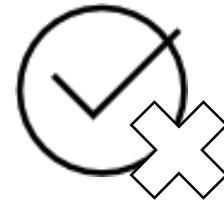
Challenges with building ledgers

RDBMS - audit tables



Hard to build

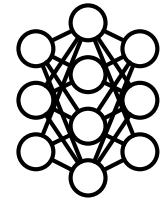
Custom audit functionality using triggers or stored procedures



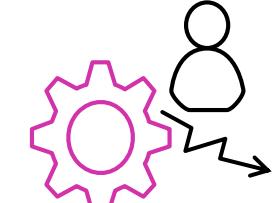
Impossible to verify

No way to verify changes made to data by sys admins

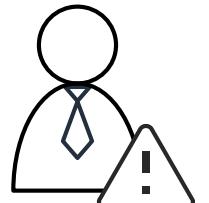
Blockchain



Adds unnecessary complexity



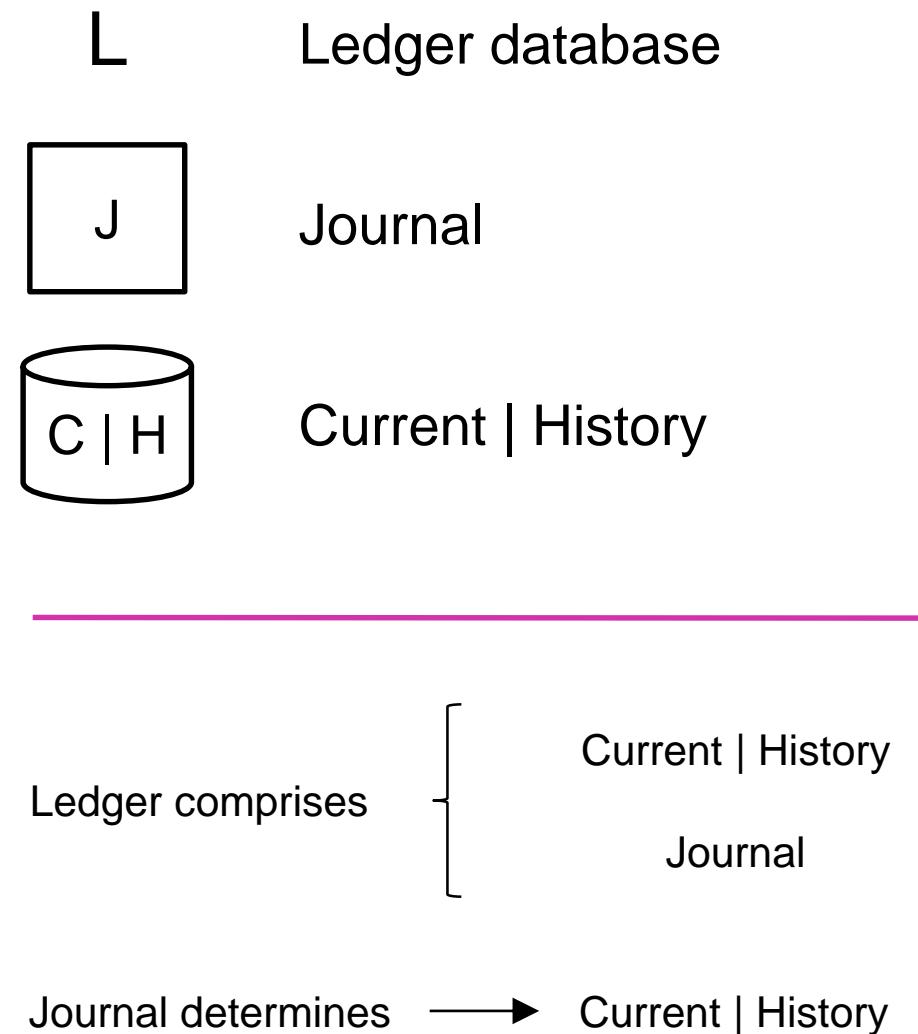
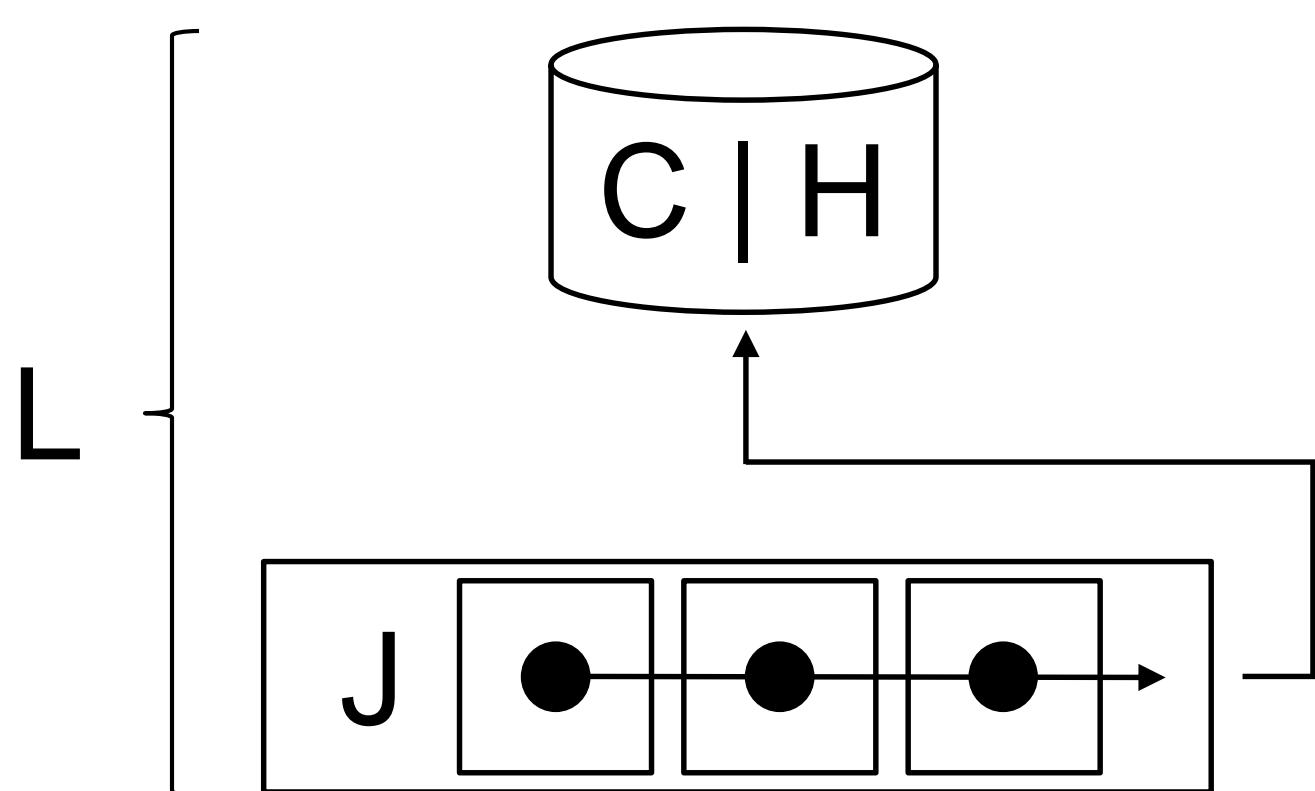
Hard to use and slow



Difficult to maintain



Ledger database concepts



How it works





How it works

C current.cars

ID	Manufacturer	Model	Year	VIN	Owner

H history.cars

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner

J

How it works



C current.cars

ID	Manufacturer	Model	Year	VIN	Owner

```
INSERT INTO cars <<  
{ 'Manufacturer':'Tesla',  
'Model':'Model S',  
'Year':'2012',  
'VIN':'123456789',  
'Owner':'Traci Russell' }  
>>
```

H history.cars

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner

J



How it works

C current.cars

ID	Manufacturer	Model	Year	VIN	Owner

```
INSERT INTO cars <<  
{ 'Manufacturer':'Tesla',  
'Model':'Model S',  
'Year':'2012',  
'VIN':'123456789',  
'Owner':'Traci Russell' }  
>>
```

H history.cars

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner

J

```
INSERT cars  
ID:1  
Manufacturer: Tesla  
Model: Model S  
Year: 2012  
VIN: 123456789  
Owner: Traci Russell
```

```
Metadata: {  
Date:07/16/2012  
}
```



How it works

C current.cars

ID	Manufacturer	Model	Year	VIN	Owner

```
INSERT INTO cars <<  
{ 'Manufacturer':'Tesla',  
'Model':'Model S',  
'Year':'2012',  
'VIN':'123456789',  
'Owner':'Traci Russell' }  
>>
```

H history.cars

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner

J

```
INSERT cars  
ID:1  
Manufacturer: Tesla  
Model: Model S  
Year: 2012  
VIN: 123456789  
Owner: Traci Russell  
  
Metadata: {  
Date:07/16/2012  
}
```



How it works

C current.cars

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Traci Russell

```
INSERT INTO cars <<  
{ 'Manufacturer':'Tesla',  
'Model':'Model S',  
'Year':'2012',  
'VIN':'123456789',  
'Owner':'Traci Russell' }  
>>
```

H history.cars

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner

J

```
INSERT cars  
ID:1  
Manufacturer: Tesla  
Model: Model S  
Year: 2012  
VIN: 123456789  
Owner: Traci Russell  
  
Metadata: {  
Date:07/16/2012  
}
```



How it works

C current.cars

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Traci Russell

```
INSERT INTO cars <<  
{ 'Manufacturer':'Tesla',  
'Model':'Model S',  
'Year':'2012',  
'VIN':'123456789',  
'Owner':'Traci Russell' }  
>>
```

H history.cars

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell

J

```
INSERT cars  
ID:1  
Manufacturer: Tesla  
Model: Model S  
Year: 2012  
VIN: 123456789  
Owner: Traci Russell  
  
Metadata: {  
Date:07/16/2012  
}
```



How it works

C current.cars

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Traci Russell

H history.cars

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell

J

INSERT cars

H (T_1)

ID:1

Manufacturer: Tesla

Model: Model S

Year: 2012

VIN: 123456789

Owner: Traci Russell

Metadata: {

Date:07/16/2012

}

How it works



C current.cars
C

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Traci Russell

```
FROM cars WHERE VIN = '123456789' UPDATE owner = 'Ronnie Nash'
```

H history.cars
H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell

J

```
INSERT cars H (T1)
ID:1
Manufacturer: Tesla
Model: Model S
Year: 2012
VIN: 123456789
Owner: Traci Russell

Metadata: {
Date:07/16/2012
}
```

How it works



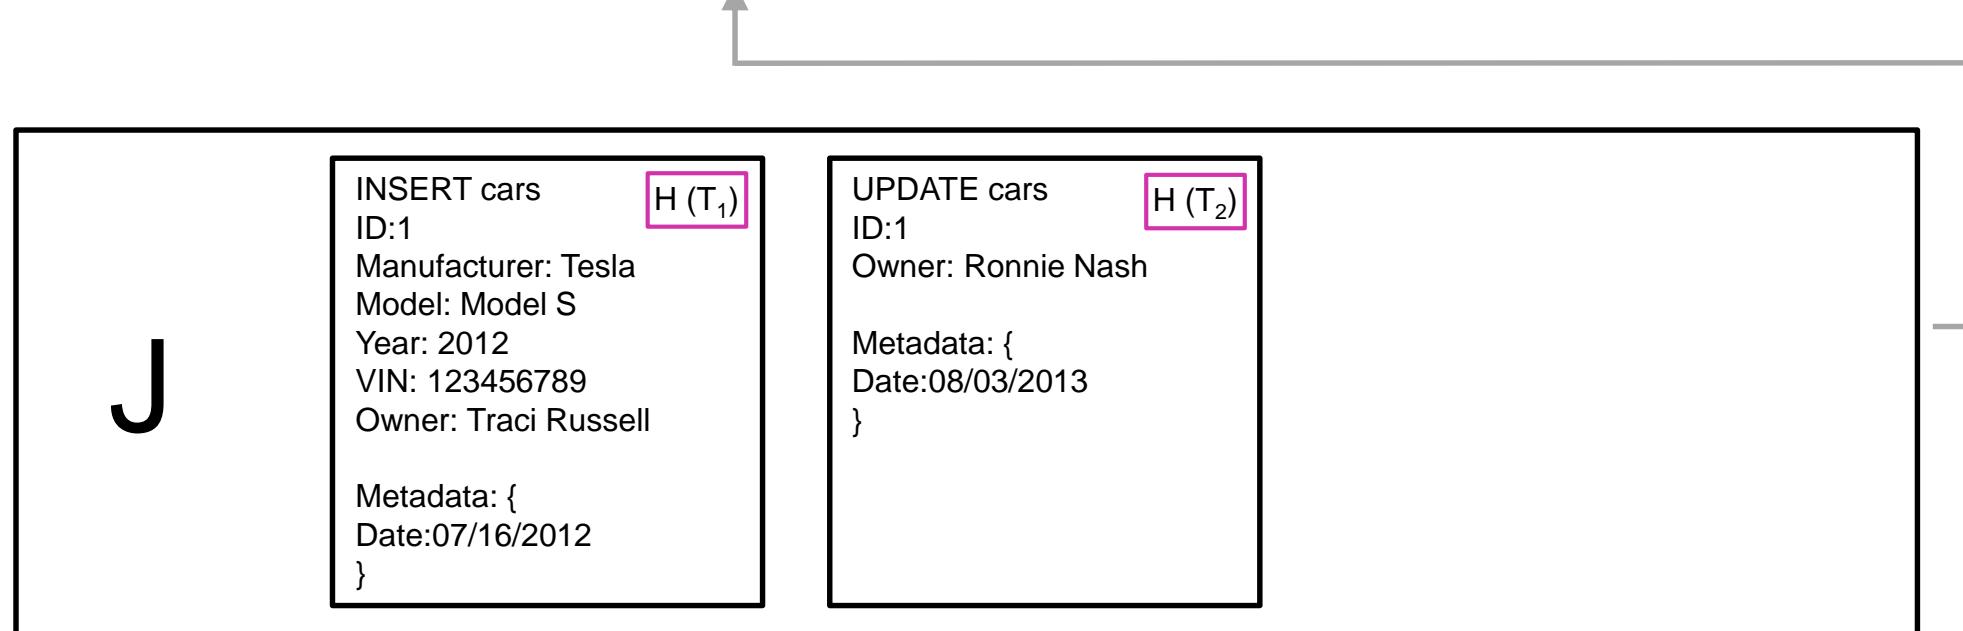
C current.cars
C

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Traci Russell

```
FROM cars WHERE VIN = '123456789' UPDATE owner = 'Ronnie Nash'
```

H history.cars
H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell



How it works



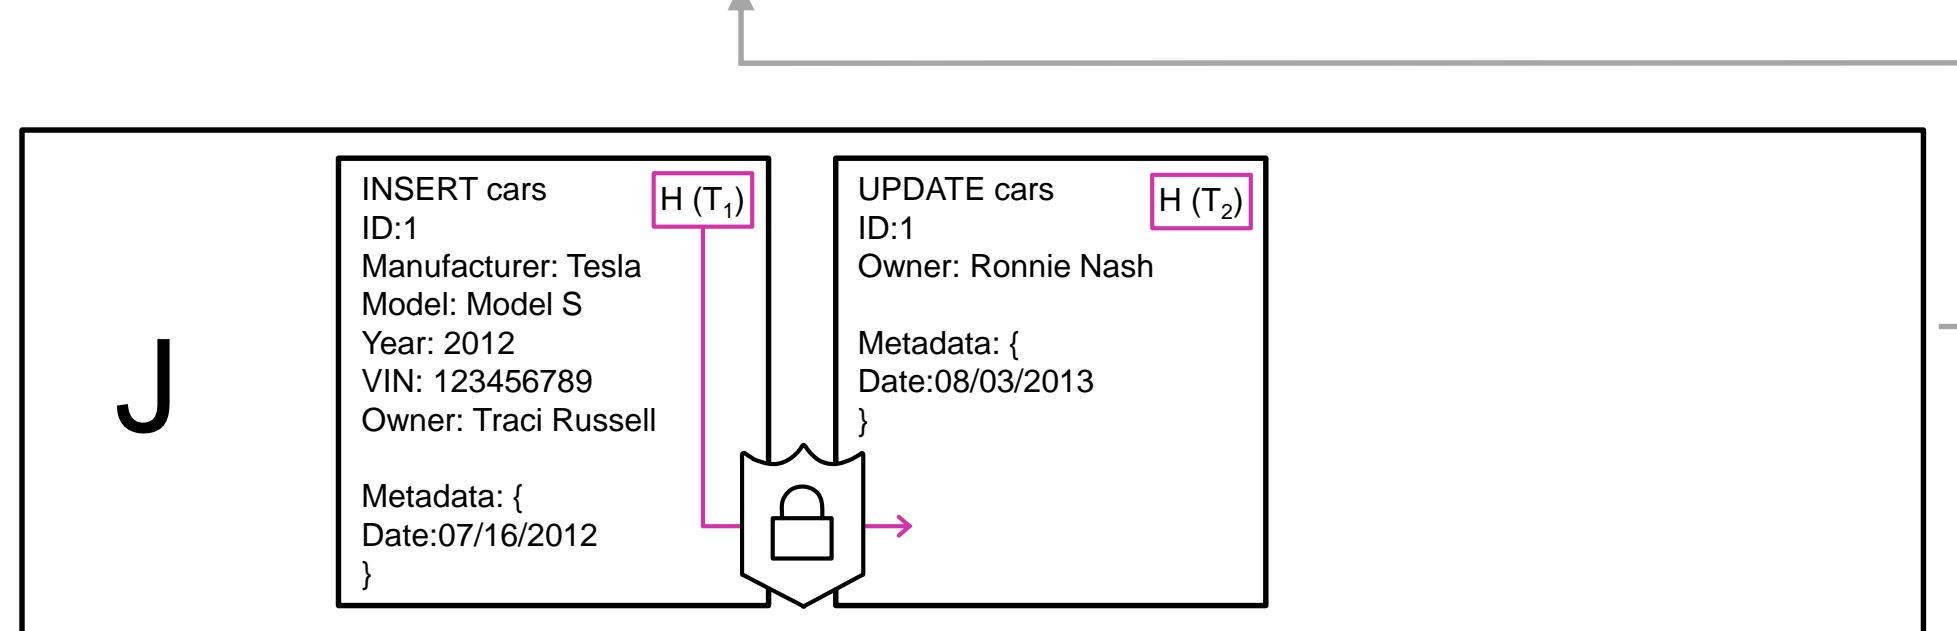
C current.cars
C

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Traci Russell

```
FROM cars WHERE VIN = '123456789' UPDATE owner = 'Ronnie Nash'
```

H history.cars
H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell



How it works



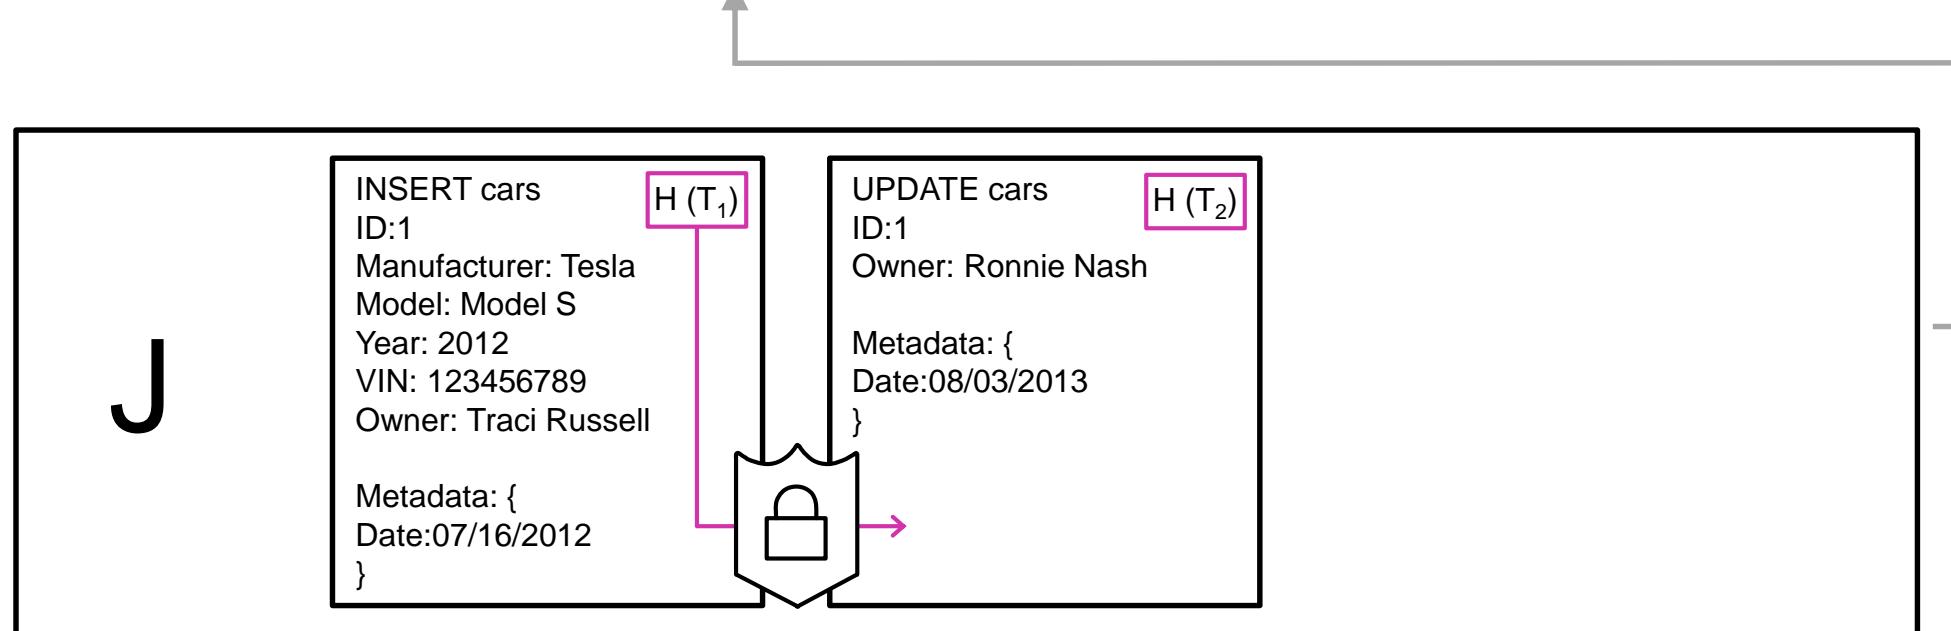
C current.cars
C

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Ronnie Nash

```
FROM cars WHERE VIN = '123456789' UPDATE owner = 'Ronnie Nash'
```

H history.cars
H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell



How it works



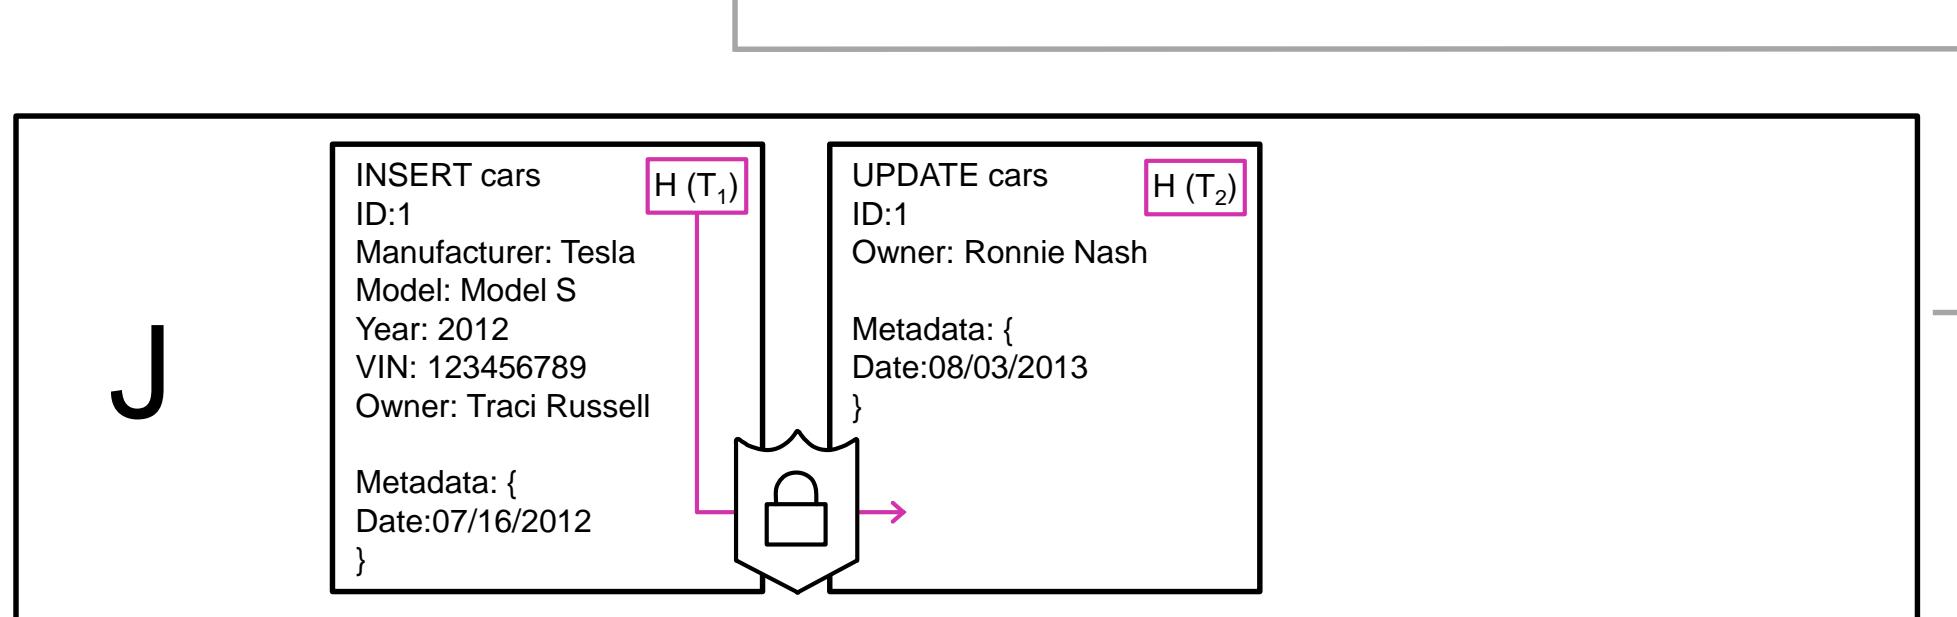
C current.cars
C

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Ronnie Nash

```
FROM cars WHERE VIN = '123456789' UPDATE owner = 'Ronnie Nash'
```

H history.cars
H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell
1	2	08/03/2013	Tesla	Model S	2012	123456789	Ronnie Nash





How it works

C current.cars

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Ronnie Nash

H history.cars

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell
1	2	08/03/2013	Tesla	Model S	2012	123456789	Ronnie Nash

J

INSERT cars
ID:1
Manufacturer: Tesla
Model: Model S
Year: 2012
VIN: 123456789
Owner: Traci Russell

Metadata: {
Date:07/16/2012
}

UPDATE cars
ID:1
Owner: Ronnie Nash

Metadata: {
Date:08/03/2013
}

How it works



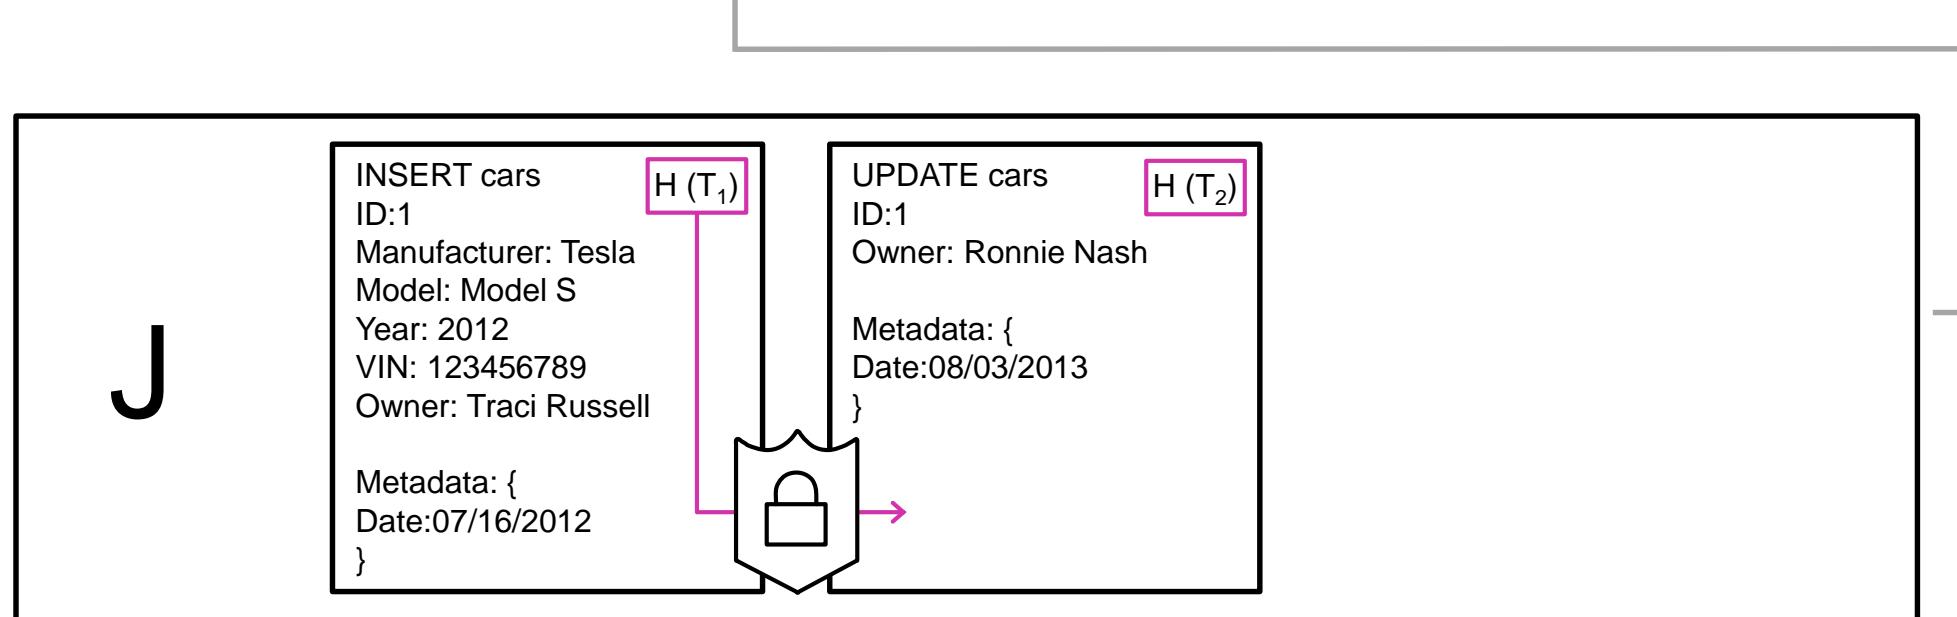
C current.cars
C

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Ronnie Nash

```
FROM cars WHERE VIN = '123456789' UPDATE owner = 'Elmer Hubbard'
```

H history.cars
H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell
1	2	08/03/2013	Tesla	Model S	2012	123456789	Ronnie Nash



How it works



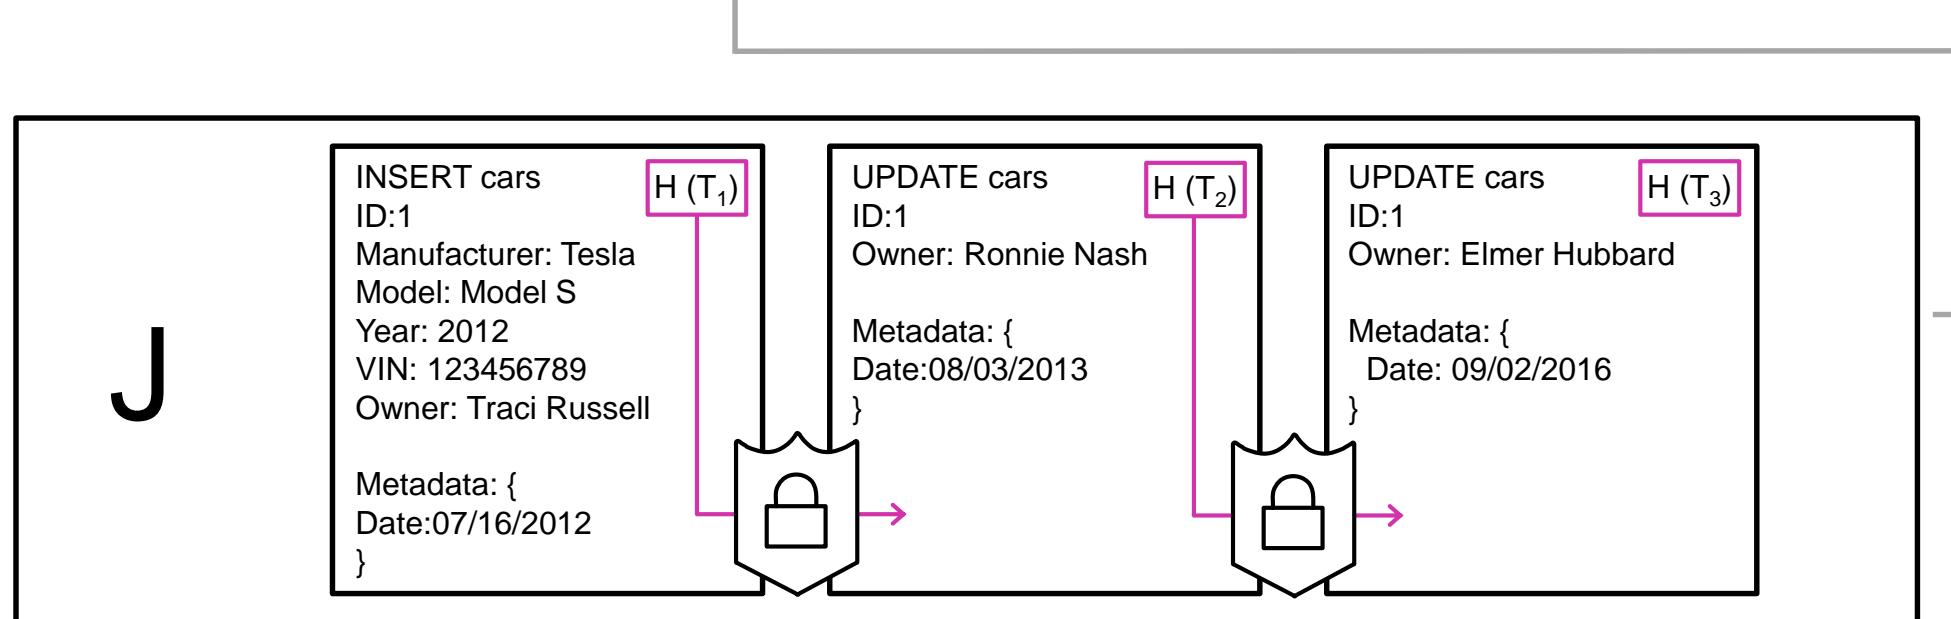
C current.cars
C

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Ronnie Nash

```
FROM cars WHERE VIN = '123456789' UPDATE owner = 'Elmer Hubbard'
```

H history.cars
H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell
1	2	08/03/2013	Tesla	Model S	2012	123456789	Ronnie Nash



How it works



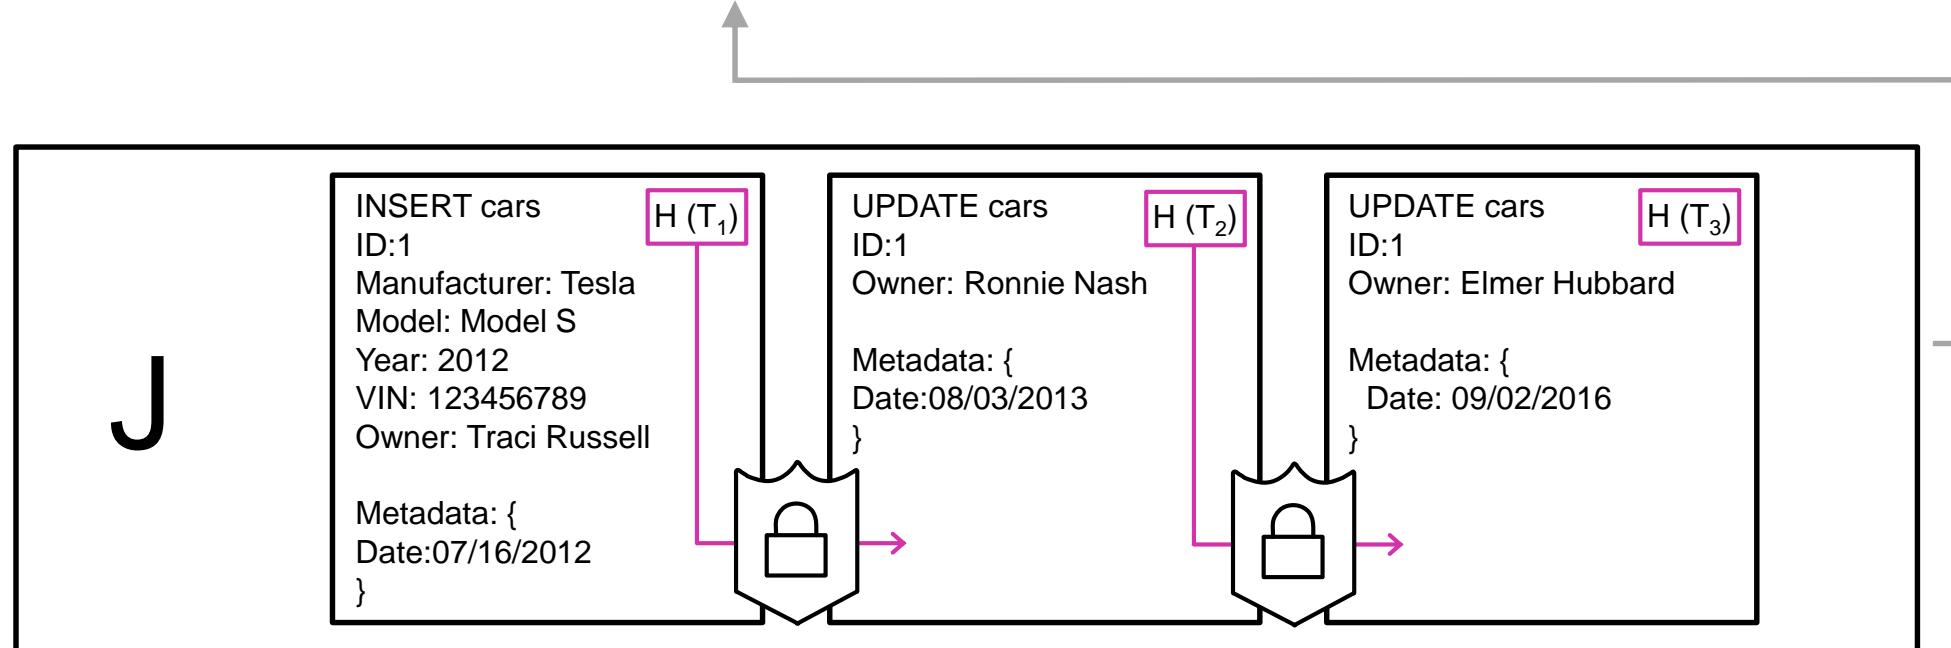
C current.cars
C

ID	Manufacturer	Model	Year	VIN	Owner
1	Tesla	Model S	2012	123456789	Elmer Hubbard

```
FROM cars WHERE VIN = '123456789' UPDATE owner = 'Elmer Hubbard'
```

H history.cars
H

ID	Version	Start	Manufacturer	Model	Year	VIN	Owner
1	1	07/16/2012	Tesla	Model S	2012	123456789	Traci Russell
1	2	08/03/2013	Tesla	Model S	2012	123456789	Ronnie Nash
1	3	09/02/2016	Tesla	Model S	2012	123456789	Elmer Hubbard





Amazon Quantum Ledger Database (QLDB)

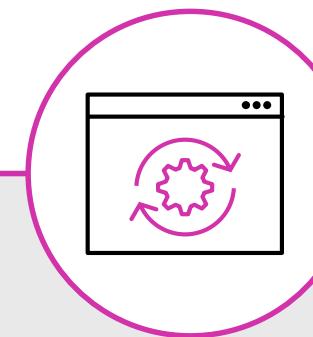
(Preview)

Fully managed ledger database

NEW!

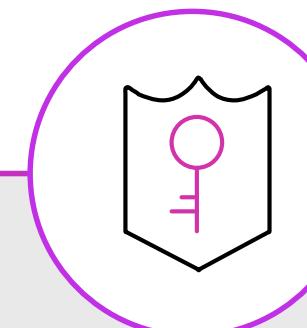
Track and verify history of all changes made to your application's data

Immutable



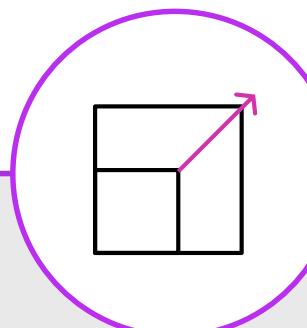
Maintains a sequenced record of all changes to your data, which cannot be deleted or modified; you have the ability to query and analyze the full history

**Cryptographically
verifiable**



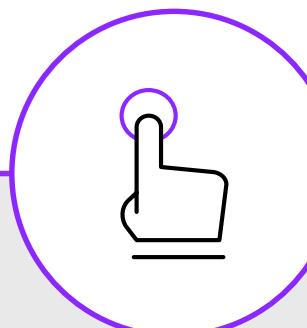
Uses cryptography to generate a secure output file of your data's history

Highly scalable



Executes 2–3X as many transactions than ledgers in common blockchain frameworks

Easy to use



Easy to use, letting you use familiar database capabilities like SQL APIs for querying the data

Time-series database



Time-series data

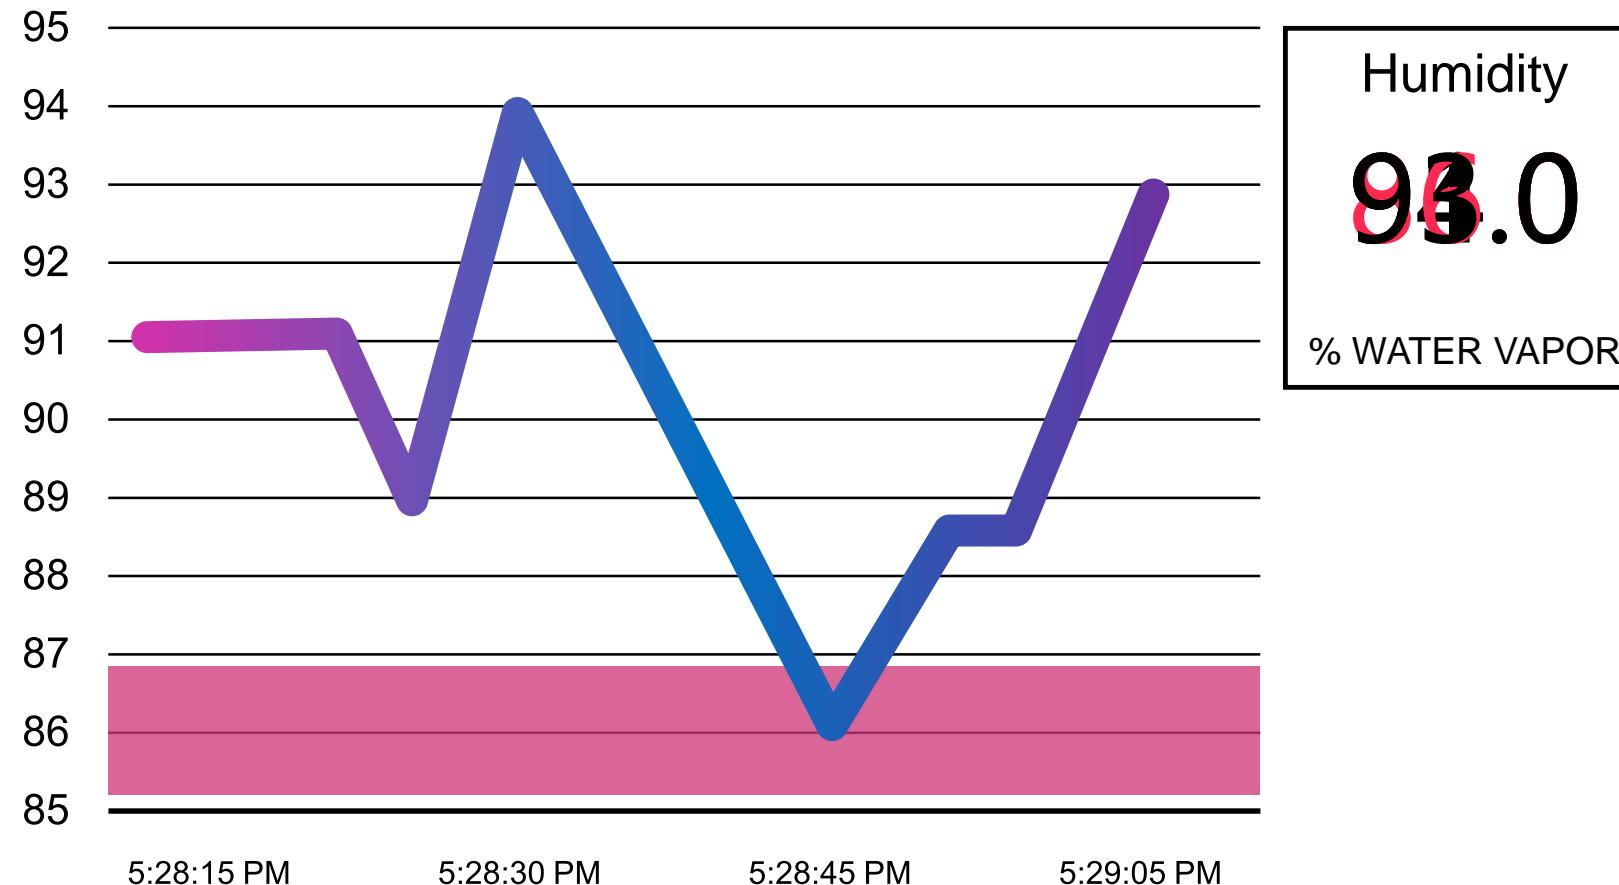
What is time-series data?

What is special about a time-series database?





Time-series use case



- ① Application events
- ② IoT Sensor Readings
- ③ DevOps data

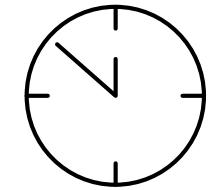


Building with time-series data is challenging

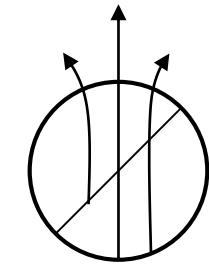
Relational databases



Unnatural for
time-series data

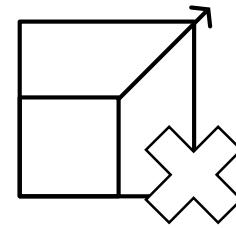


Inefficient
time-series data
processing

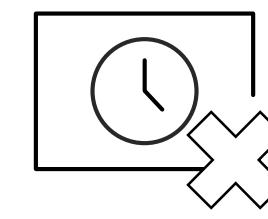


Rigid schema
inflexible for fast
moving time-series
data

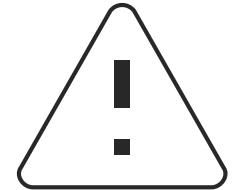
Existing time-series databases



Difficult to scale



Difficult to
maintain high
availability



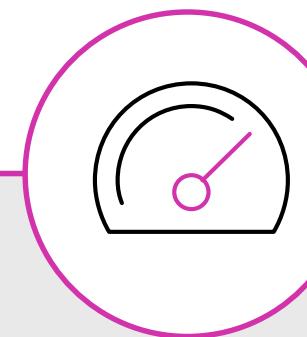
Limited data
lifecycle
management

Amazon Timestream (sign up for the prev NEW!)



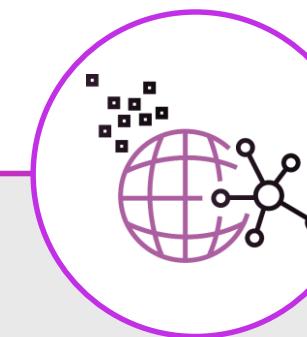
Fast, scalable, fully managed time-series database

1,000x faster and 1/10th the cost of relational databases



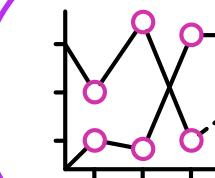
Collect data at the rate of millions of inserts per second (10M/second)

Trillions of daily events



Adaptive query processing engine maintains steady, predictable performance

Time-series analytics



Built-in functions for interpolation, smoothing, and approximation

Serverless



Automated setup, configuration, server provisioning, software patching

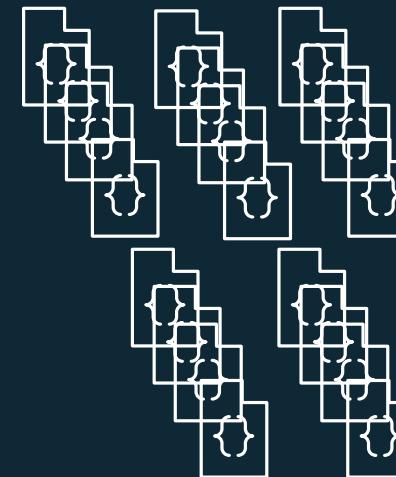
Document database

What is a document database?

Non-relational database that stores semi-structured data as JSON documents

Flexible schema – each document can contain different fields

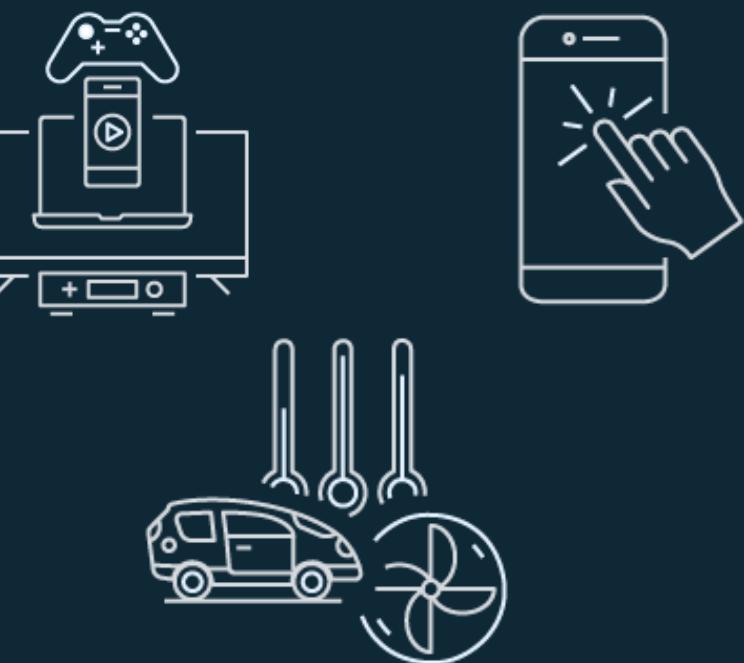
Semi-structured
data sets



Highly scalable



Modern applications



Why use a document database?

The JSON document model maps naturally to application data



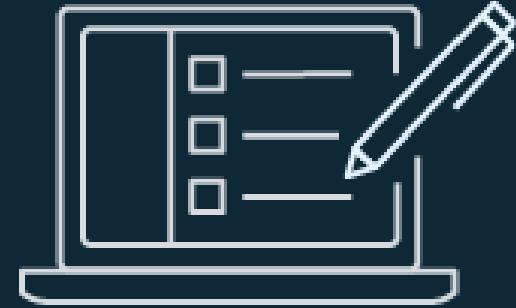
Each document can have a different data structure and is independent of other documents



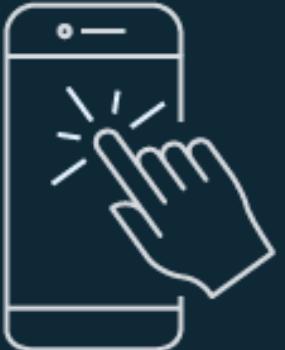
Index on any key in a document, and run ad hoc and aggregation queries across your data set



Use cases for document databases



Content
Management



Mobile



Personalization



Catalog



Retail and
Marketing



User profiles

MongoDB – popular document database engine

Open-source

Over 40 million downloads

Easy for developers to get started

MongoDB API is rich and powerful

Easy to programmatically query

Rank			DBMS
Jan 2019	Dec 2018	Jan 2018	
1.	1.	1.	Oracle 
2.	2.	2.	MySQL 
3.	3.	3.	Microsoft SQL Server 
4.	4.	4.	PostgreSQL 
5.	5.	5.	MongoDB 
6.	6.	6.	IBM Db2 
7.	7.	↑ 9.	Redis 
8.	8.	↑ 10.	Elasticsearch 
9.	9.	↓ 7.	Microsoft Access
10.	10.	↑ 11.	SQLite 

Source: <http://db-engines.com>

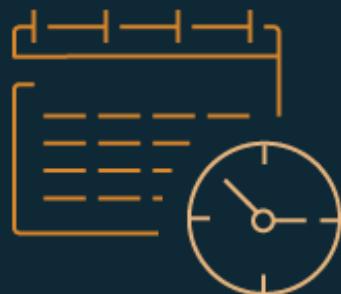
Running MongoDB is difficult.....



Expensive



Experts only



Time to market



What if you could create and scale
MongoDB compatible database
clusters in minutes?

Introducing:



Amazon
DocumentDB

Fast

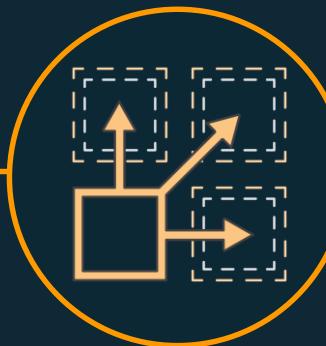
Fast, reliable, and fully-managed MongoDB-compatible database service

Fast



Millions of requests per second with millisecond latency; scale-out up to 15 read replicas

More throughput



Separation of storage and compute offloads replication, providing 2x the throughput of current MongoDB managed services

Automatic scaling



DocumentDB will automatically grow the size of your storage volume as your cluster storage needs grow.

Analytics



Scale-up instances in minutes for analytical queries and scale them down at the end of the day

Reliable

Fast, **reliable**, and fully-managed MongoDB-compatible database service

Automatic failure recovery



Failing instances are automatically detected and recovered; no cache warm-up needed

Automatic failover



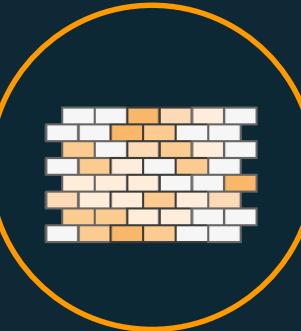
Replicas are automatically promoted to primary

Point-in-time restore



Automated backups are stored in Amazon S3, designed for 99.999999999% durability

Durable



Data is replicated six-ways across three AZs

Fully-managed

Fast, reliable, and **fully-managed** MongoDB-compatible database service

Automatic patching



Up-to-date with the latest patches

Quick start



Provision production-ready clusters in minutes

Monitoring



Over 20 key operational metrics for your clusters at no extra charge

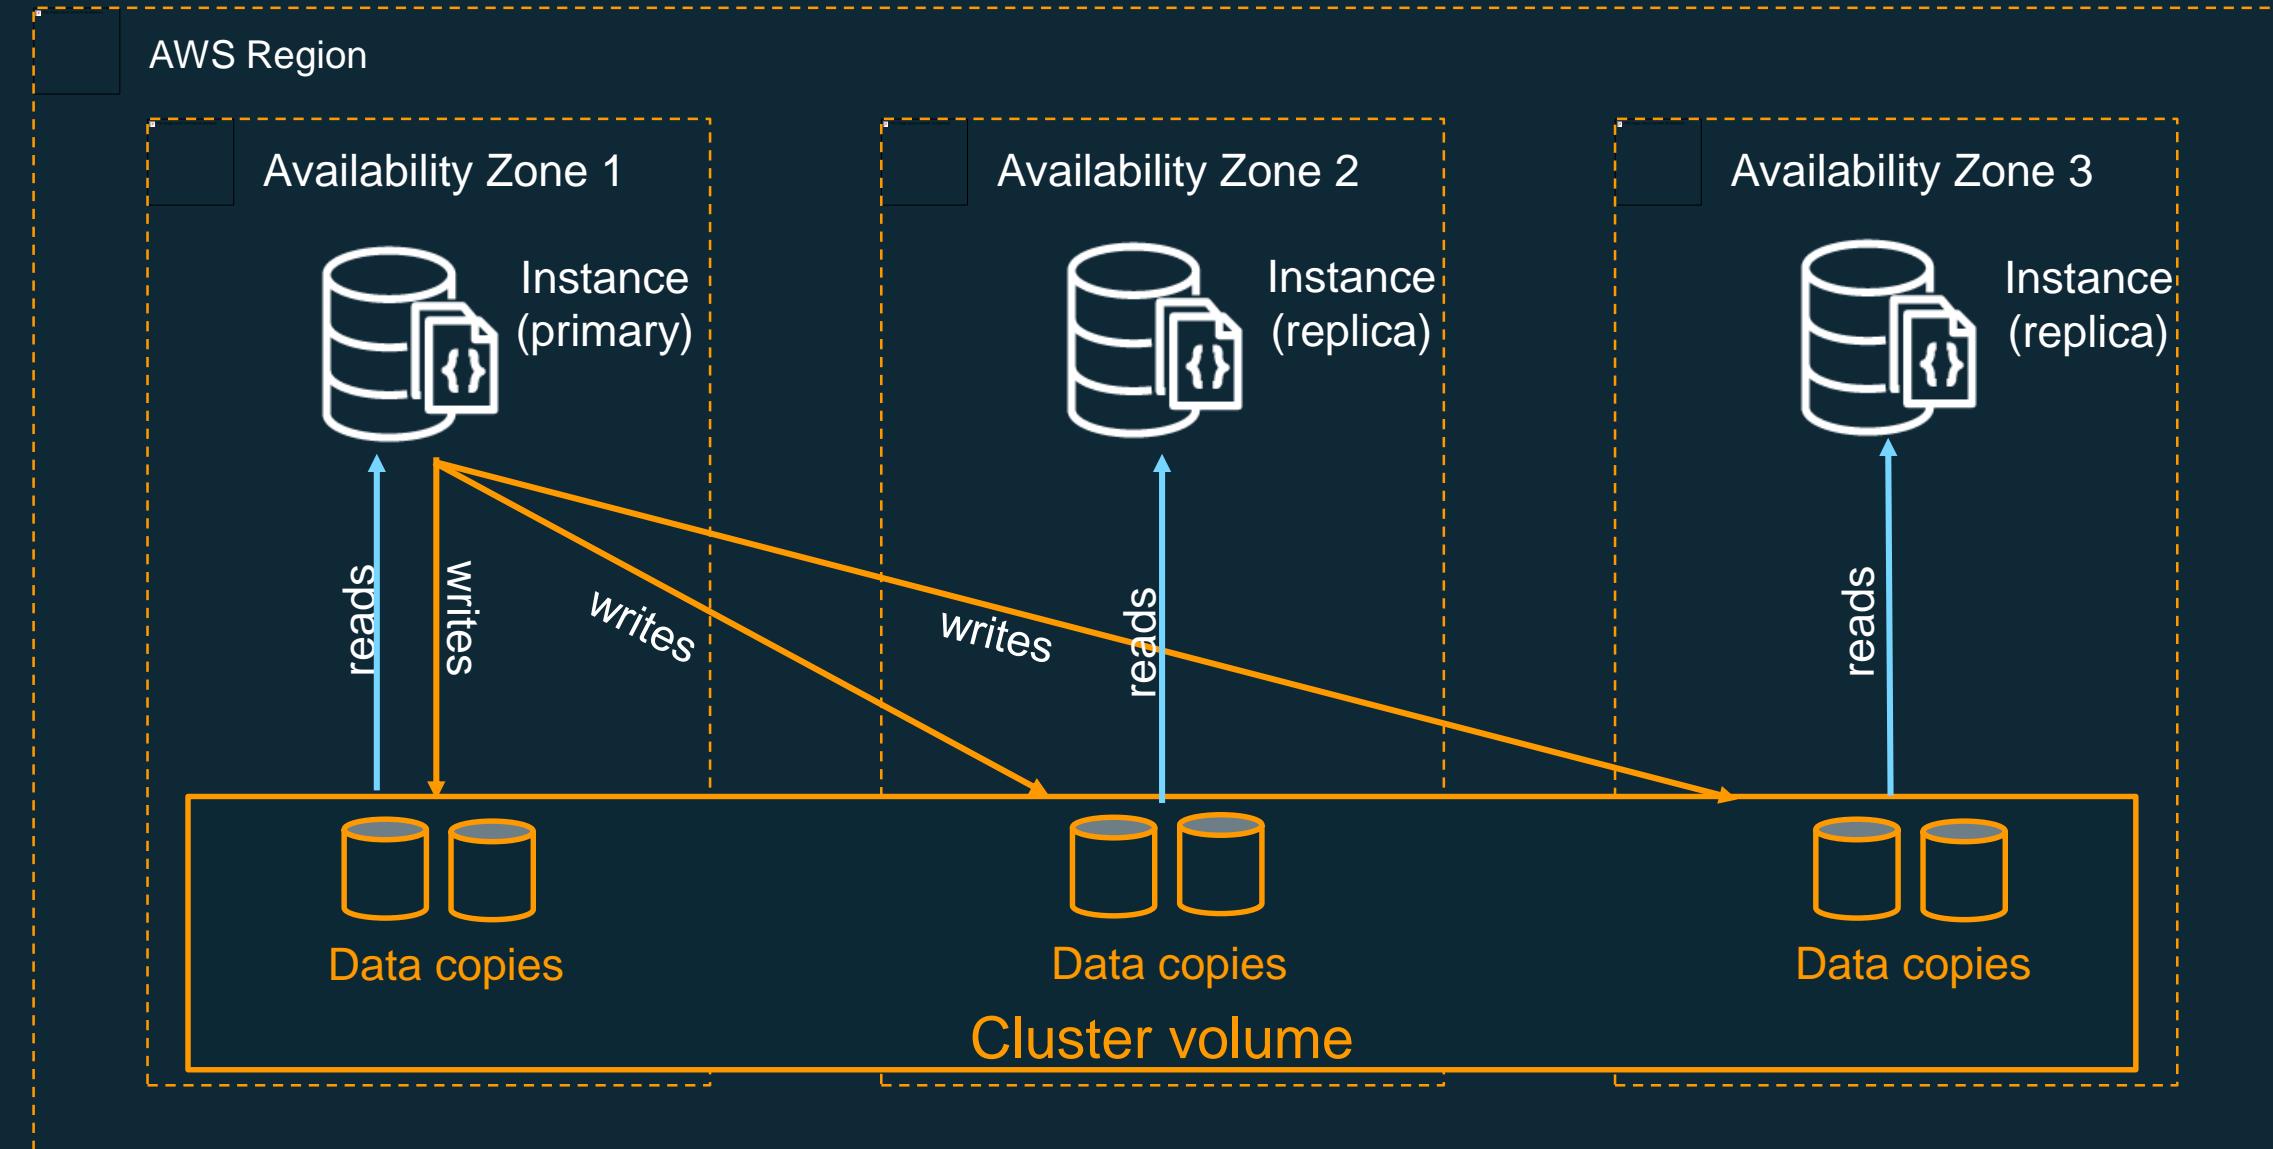
Integrated



Deeply integrated with AWS services such as CloudFormation, CloudTrail, CloudWatch, DMS, IAM, VPC, and more.

DocumentDB Architecture

Separate compute and storage provide 2x throughput of current MongoDB managed services



MongoDB-compatible

Fast, reliable, fully-managed MongoDB-compatible database

MongoDB 3.6



Compatible with MongoDB Community Edition 3.6

Same drivers, tools



Use the same MongoDB drivers and tools with DocumentDB; as simple as changing an application connection string

Migration with DMS



Live migrations with DMS; free for 6-months

Replica sets



Read scaling is easy with automatic replica set configurations

Amazon DocumentDB availability

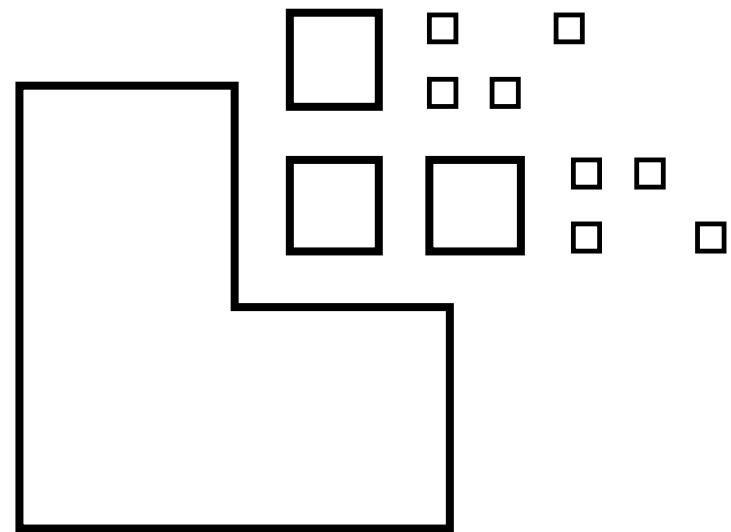


- US West (Oregon)
- US East (N. Virginia)
- US East (Ohio)
- EU (Ireland)

Summary

One size doesn't fit all



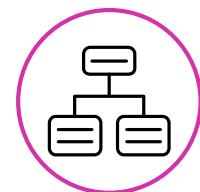


Purpose-built databases

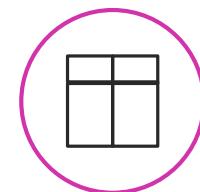
The right tool for the right job



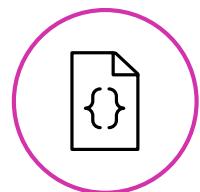
Purpose-built databases



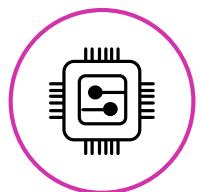
Relational



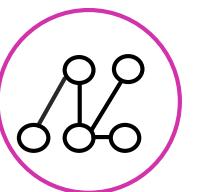
Key-value



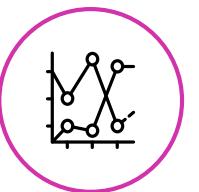
Document



In-memory



Graph



Time-series



Ledger



**Amazon
RDS**

Aurora

Community

Commercial



ORACLE®



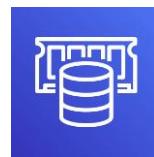
Microsoft®
SQL Server®



DynamoDB



DocumentDB



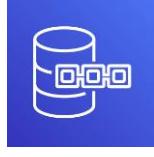
ElastiCache



Neptune



Timestream



QLDB

Redis Memcached

Q & A

Thank you

vladsim@amazon.com