



Cost-optimize your workloads on Amazon EKS

Borja Pérez Guasch bpguasch@amazon.es

Specialist Solutions Architect, Flexible Compute



Agenda

EC2 Spot & allocation strategies

Graviton

Amazon EKS & Karpenter

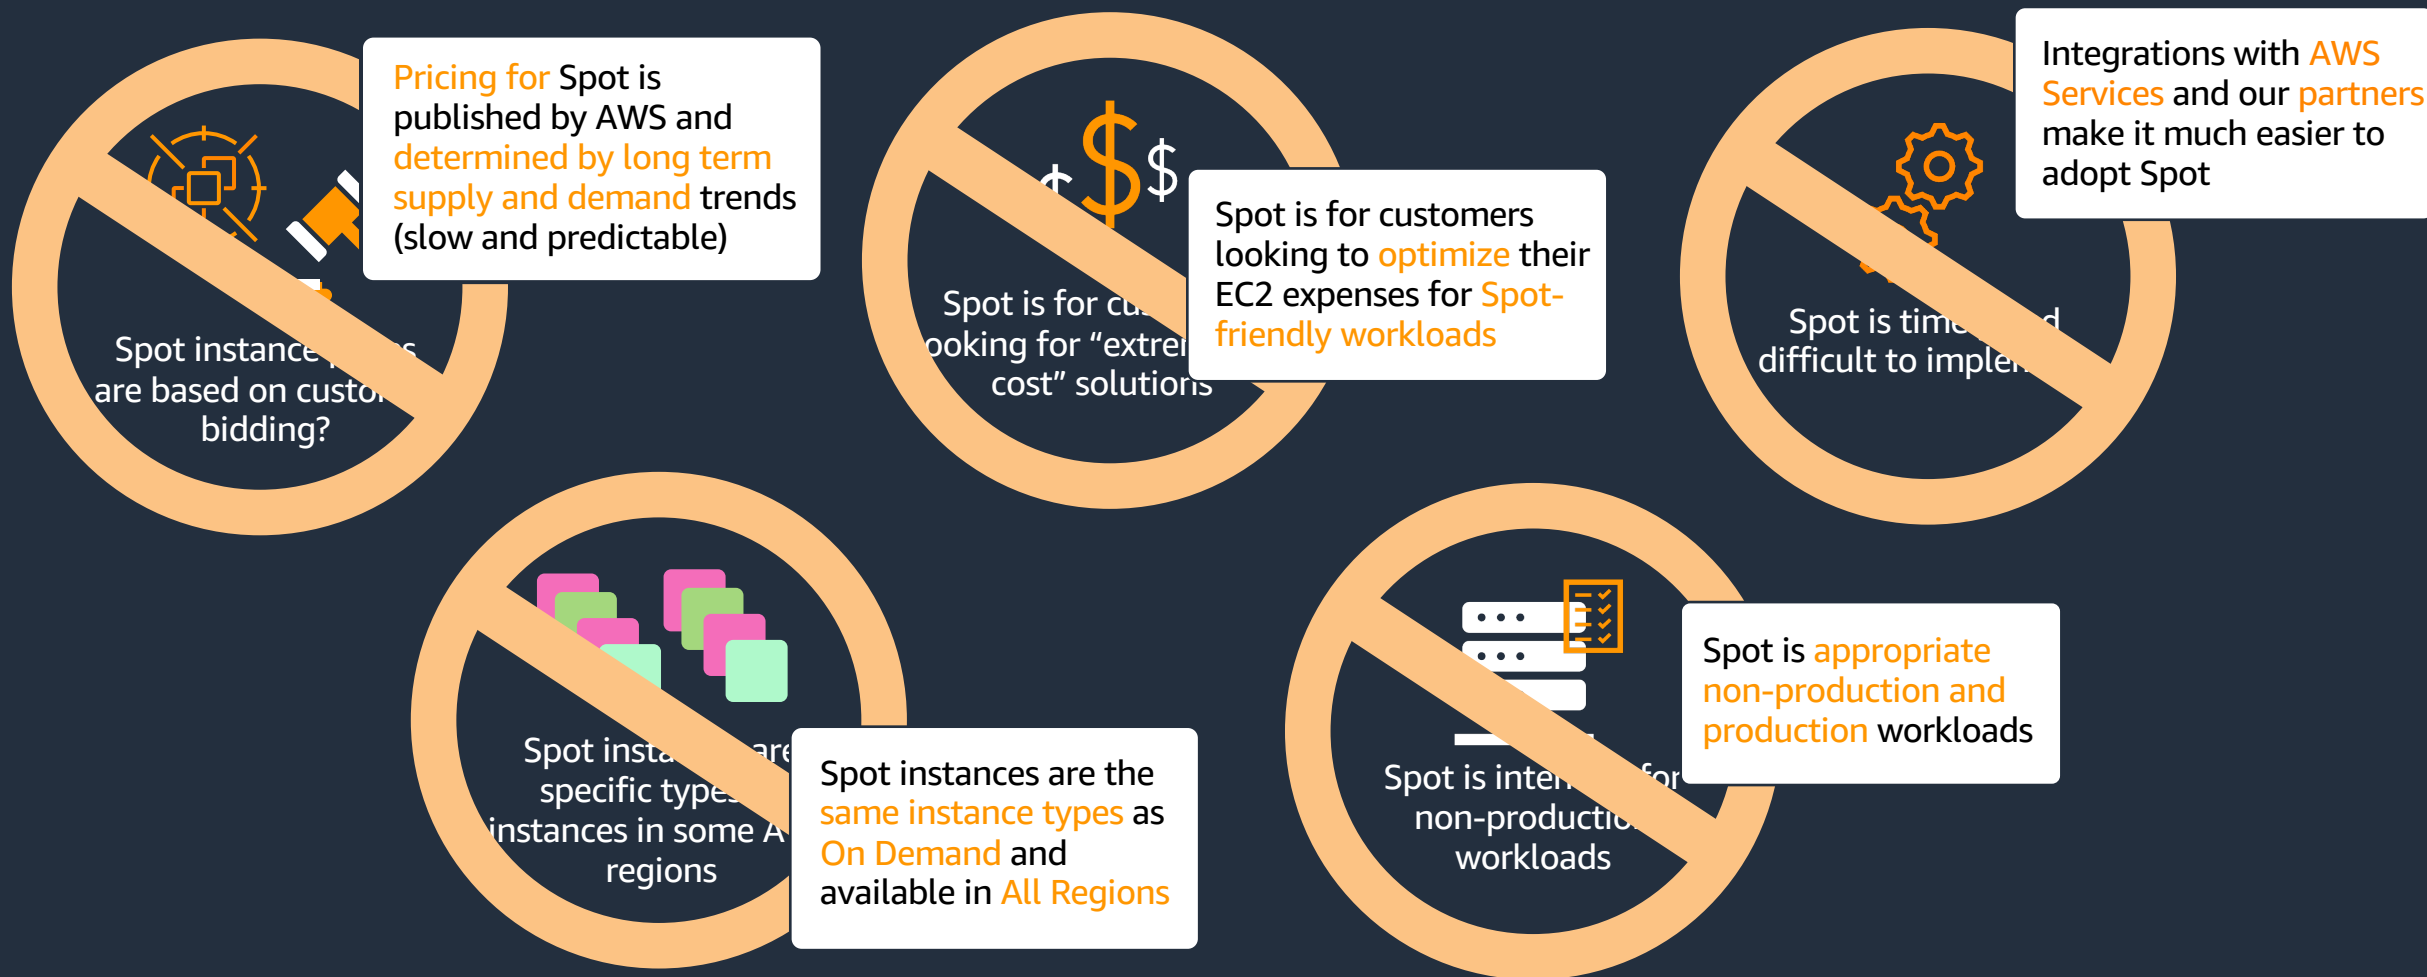
Kubecost


Q&A

EC2 Spot & allocation strategies



EC2 Spot knowledge check






MORE THAN
600
INSTANCE TYPES

The picture represents an **example** of how a particular instance type-size could be in use and have spare capacity


The white squares here represent the spare capacity of an instance type-size and the group of spare instances we call **Spare Capacity Pools**.




Customers can provision spare instances at a discounted rate – we call these **EC2 Spot instances**.

EC2 Spot instances offer up to a **90%** discount compared to on-demand.

How Spot Instances Work ...



Spot instances Spare EC2 Capacity



Provisioned from spare-capacity, uses same infrastructure as On-Demand (OD)

Pricing Up to 90% off compared to On-Demand




Not bidding - Pricing is based on long-term supply and demand, smooth and predictable

Interruption only interrupted if OD needs capacity




AWS can reclaim with 2-minute notice; issues two types of notifications to help handle interruptions

Diversification and flexibility is key



Make use of different instance types, sizes, Availability Zones, and times Regions

Interruptions



By the nature of Spot as spare-capacity, a Spot instance can be interrupted if the instance is needed by On-Demand.

AWS provides two types of notifications to enable you to handle the response in an **automated** way:

EC2 instance rebalance recommendation (proactive)



- Spot instance is at elevated risk of interruption
- Built in support for AWS integrations such as **EC2 Auto Scaling** and **EKS Managed Node Groups**



Spot instance termination notice (reactive)




- Interruption of instance will happen in 2 minutes, adjust your workload appropriately
- Built in support for AWS integrations such as **EC2 Auto Scaling** and **EKS Managed Node Groups**



Historically 95% of the Spot instances launched in the last 3 months completed without interruption

Diversification across instances reduces interruptions



Diversification + Flexibility is key...

1. Instance Flexible (Type + Size!)

- Use as many instance types as possible that suits the workload
- Multiple instance types are key to resilient clusters
- Attribute based instance selection helps you choose ranges of instance types

2. Availability Zone Flexible

- Capacity exists differently across availability zones, and multiplies potential capacity based on how many AZs are used

3. Region Flexible

- Different capacity exists across regions
- HPC customers and high production Spot users may span regions

4. Time Flexible

- Capacity can differ based on time/region usage, it is sometimes worth exploring running workloads at different times to utilize spare capacity



Use the right EC2 Spot allocation strategy




On-demand allocation

- Prioritized
- Lowest-price (default for mixed instance groups)

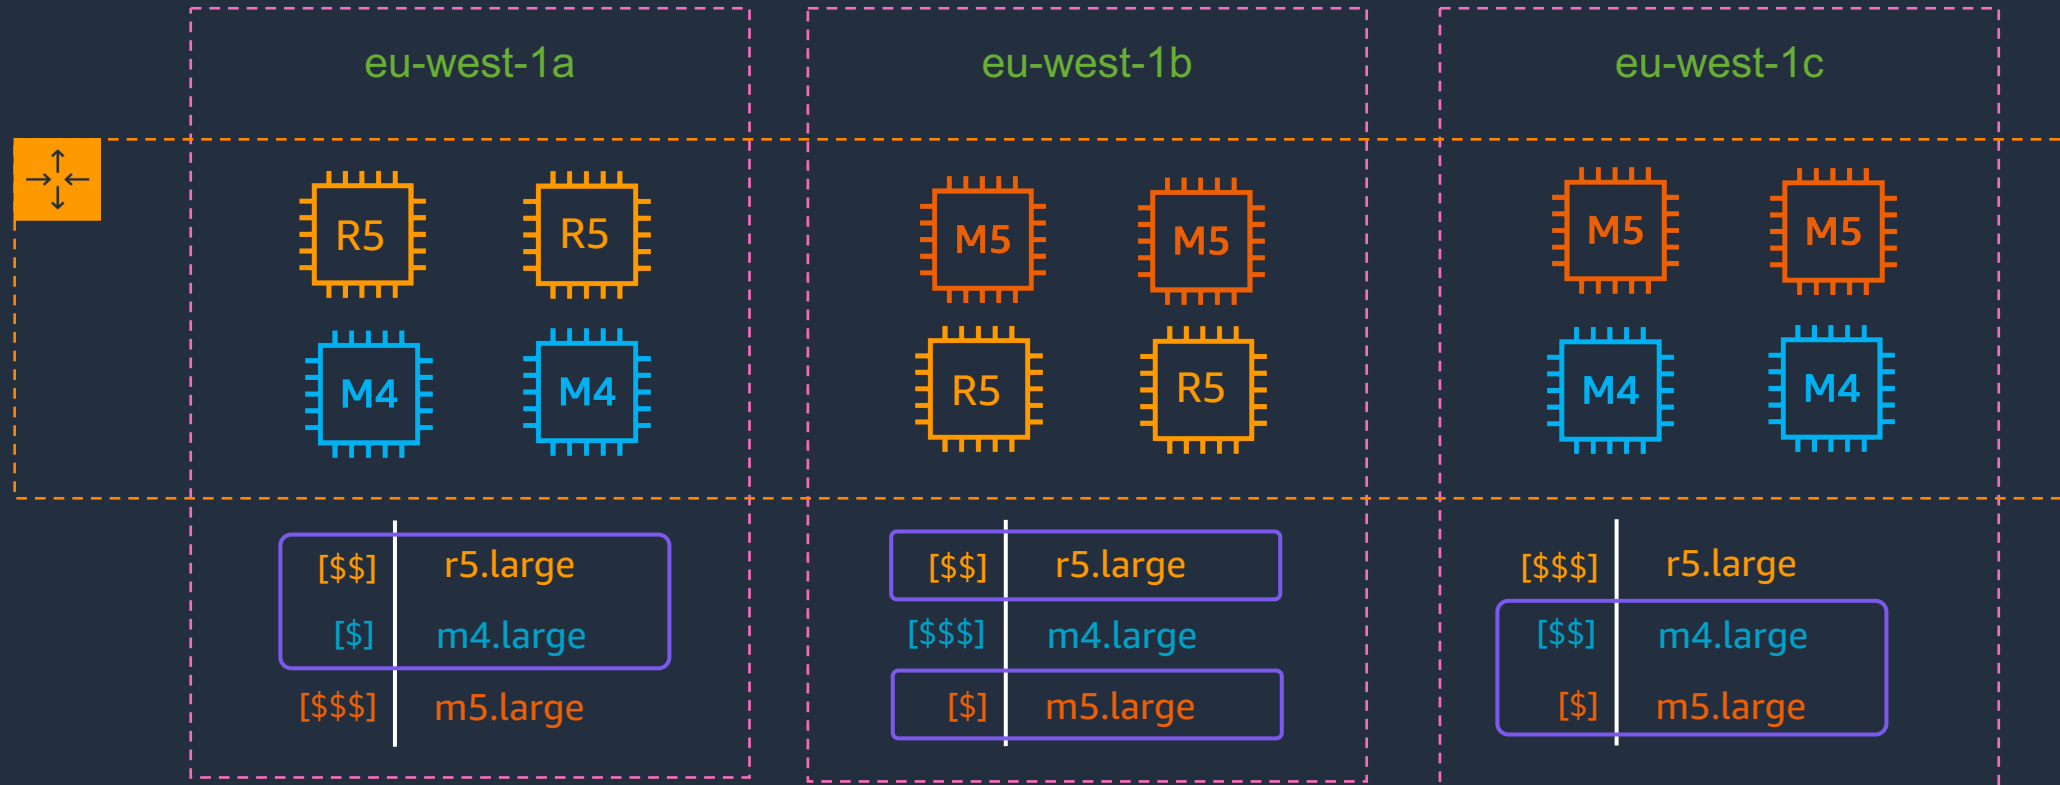


Spot allocation

- Lowest-price
- Capacity-optimized
- Price-capacity-optimized (recommended) 



Allocation Strategy: lowest-price

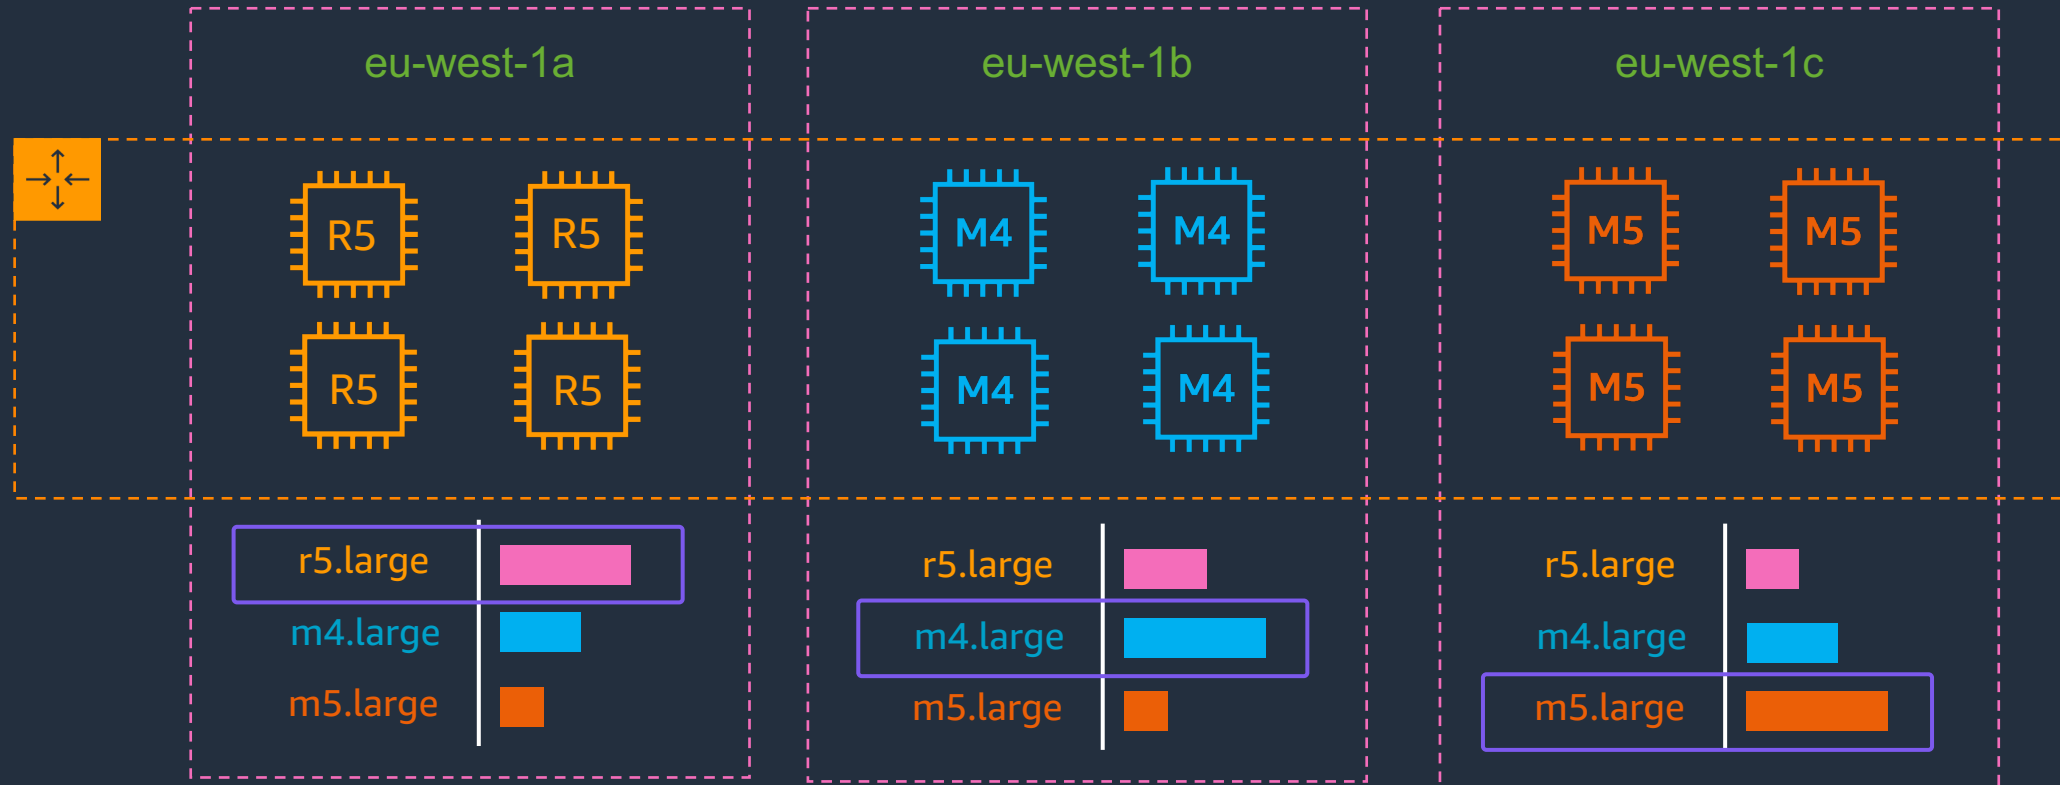


spotAllocationStrategy: **lowest-price** SpotInstancePools:2 (default)

overrides: ["r5.large", "m4.large", "m5.large"]

Desired capacity: **12** OnDemandBaseCapacity: 0 OnDemandPercentageAboveCapacity: 0

Allocation Strategy: capacity-optimized



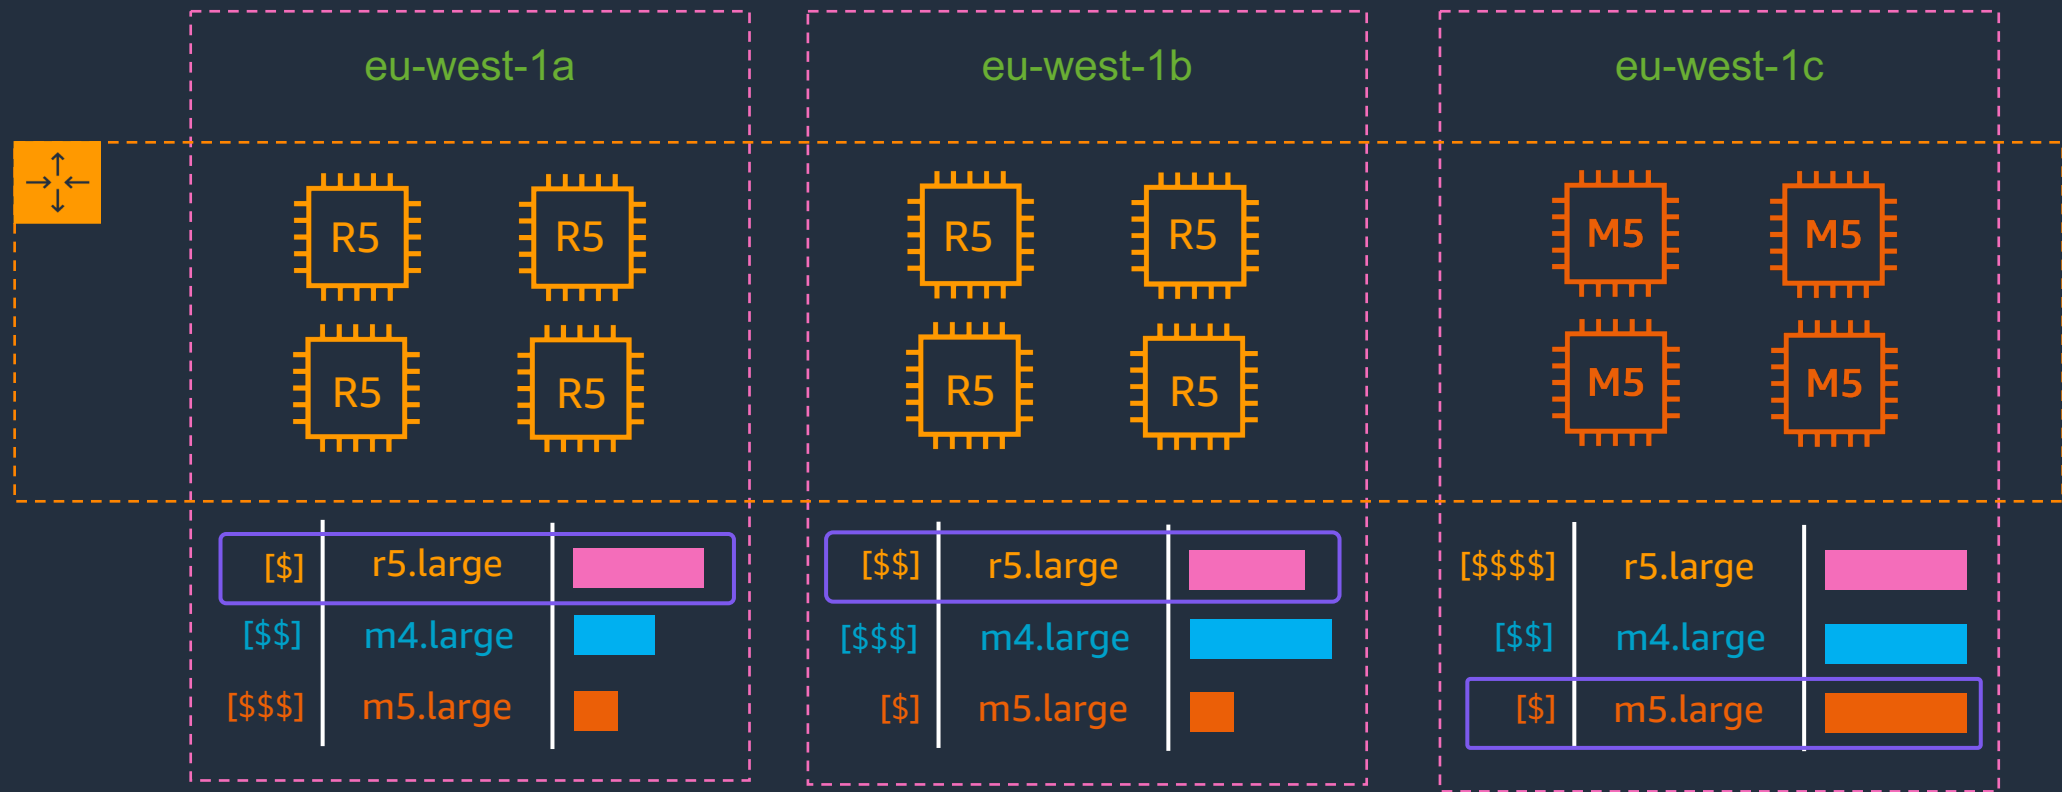
spotAllocationStrategy: **capacity-optimized**

overrides: ["r5.large", "m4.large", "m5.large"]

Desired capacity: **12** OnDemandBaseCapacity: 0 OnDemandPercentageAboveCapacity: 0



Allocation Strategy: price-capacity-optimized



spotAllocationStrategy: price-capacity-optimized

overrides: ["r5.large", "m4.large", "m5.large"]


Desired capacity: 12 OnDemandBaseCapacity: 0 OnDemandPercentageAboveCapacity: 0

Adapt your EC2 Purchasing Strategy to your Workload

USE **RIs** AND **SAVINGS PLANS**
FOR KNOWN/
STEADY-STATE WORKLOADS

SCALE USING **ON-DEMAND**
FOR NEW OR STATEFUL
SPIKY WORKLOADS

SCALE USING **SPOT INSTANCES**
FOR FLEXIBLE, FAULT-
TOLERANT WORKLOADS



AWS Graviton



Broadest choice of processors



Intel® Xeon
Scalable
processors



AMD EPYC
processors



AWS Graviton
processors



Apple M1
processors

x86

Arm64

AWS Graviton



Up to **40% better price-performance** for a broad spectrum of workloads




Up to **20% less expensive** than comparable x86-based instances



Up to **60% more energy efficient** vs. comparable x86-based instances

Graviton3 and Amazon EC2 C7g instances



Up to 25% better performance compared to Graviton2

Up to 2x higher floating-point performance, up to 2x faster cryptographic workload performance, and up to 3x better machine learning performance compared to Graviton2

First in the cloud to feature DDR5 memory

60% more energy efficient over comparable EC2 instances

C7g instances provide the best price performance for compute-intensive workloads in Amazon EC2

AWS Graviton: Broad workload applicability


Web and gaming servers




Open-source databases




High performance computing




In-memory caches



Media encoding



Electronic design automation



Analytics




Microservices




AWS managed services supporting Graviton


Databases




Amazon
DocumentDB



Amazon
Aurora



Amazon
RDS



Amazon
ElastiCache




Amazon
MemoryDB




Amazon
Neptune

Analytics



Amazon
OpenSearch




Amazon
EMR


Compute



AWS
Lambda



AWS Fargate
(for ECS)



AWS Elastic
Beanstalk

Machine Learning



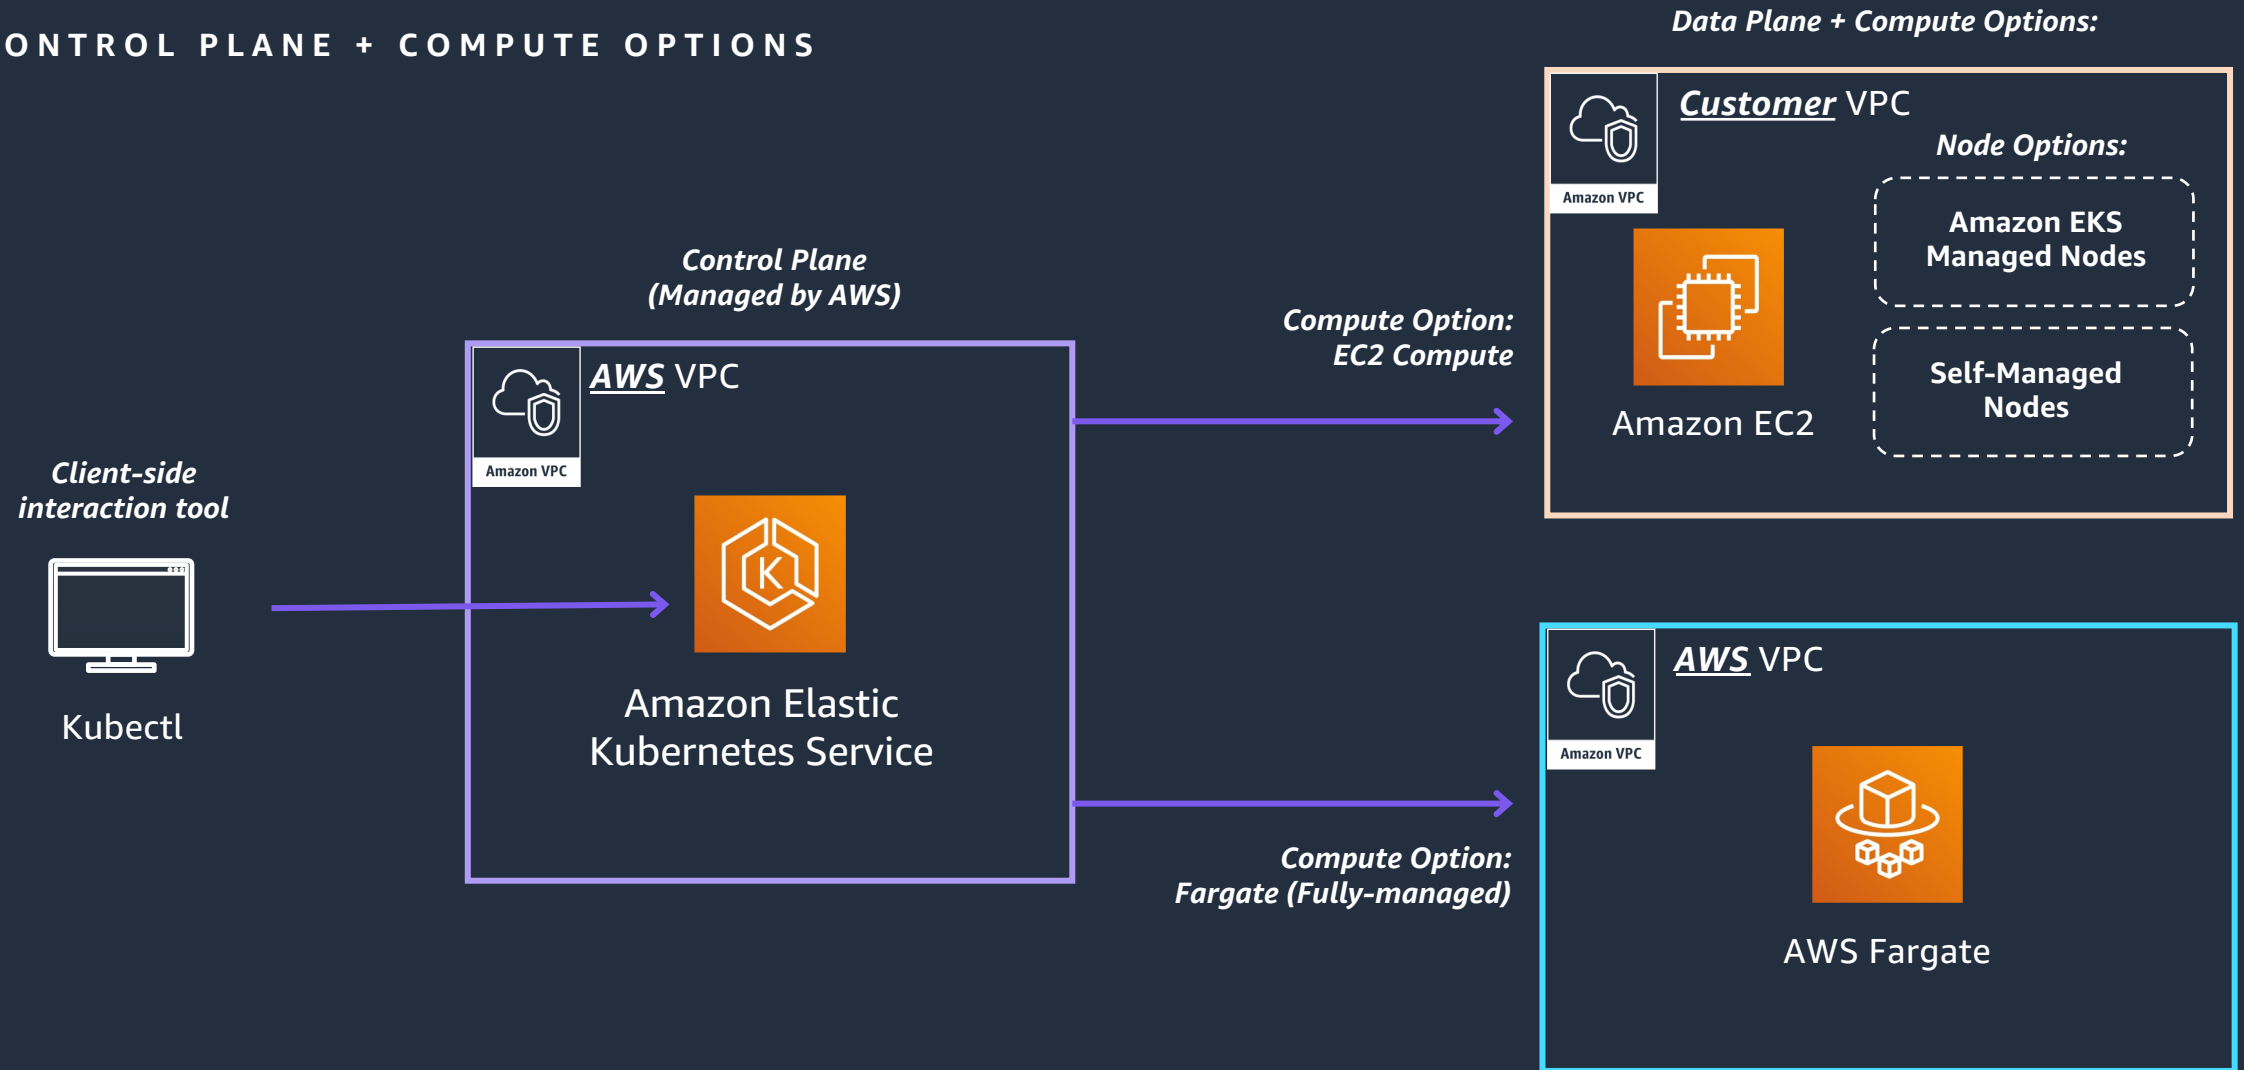
Amazon
SageMaker

Amazon EKS & Karpenter

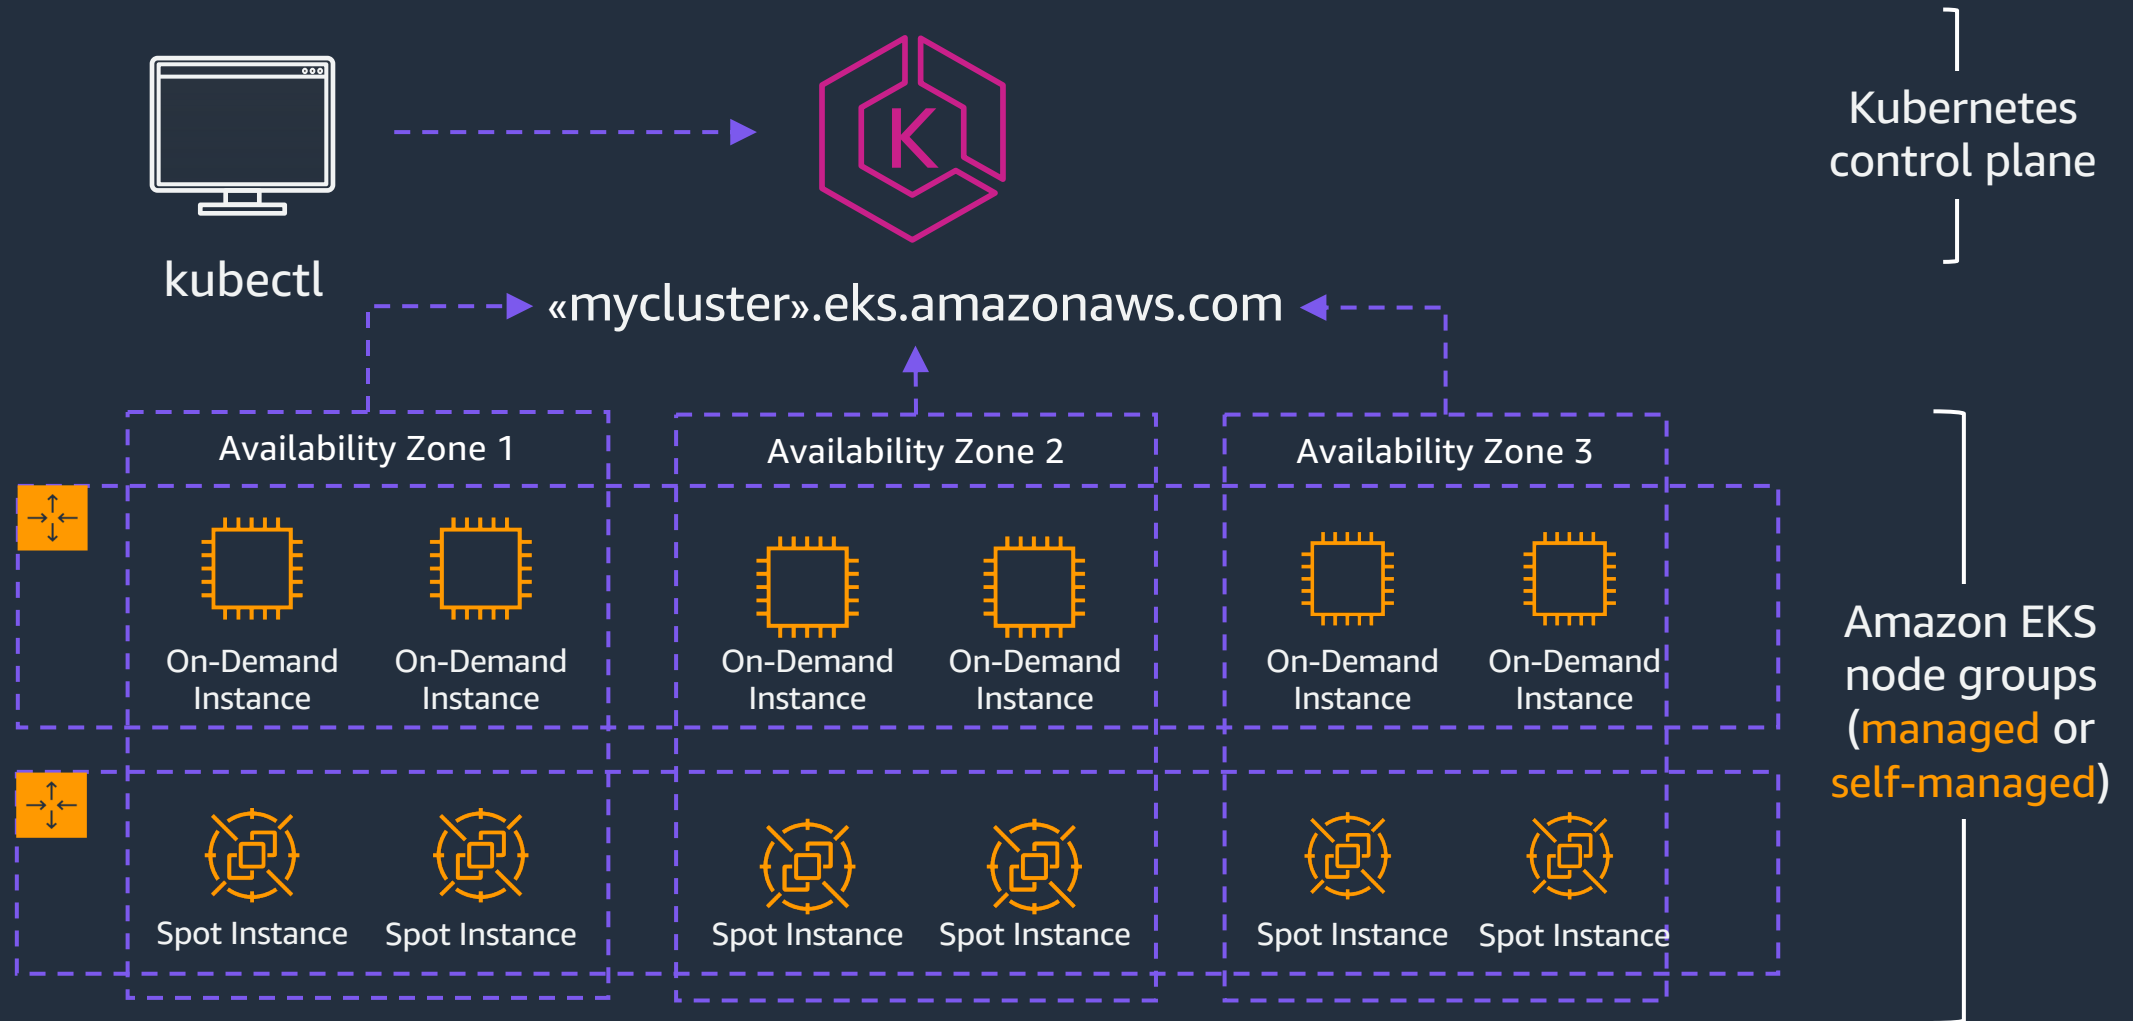


Amazon Elastic Kubernetes Service

CONTROL PLANE + COMPUTE OPTIONS




EKS Cluster Architecture and Node Provisioning



Karpenter


EKS CLUSTER NODE PROVISIONER



Application First Infrastructure

Node provisioning based on Pod requirements


Default infrastructure provisioning



Simplified Configuration

Single configuration with On-demand and Spot purchasing options and diverse instance types

Track nodes using native Kubernetes labels



Diversify across Spot and On-Demand

Simplified diversification across purchase options



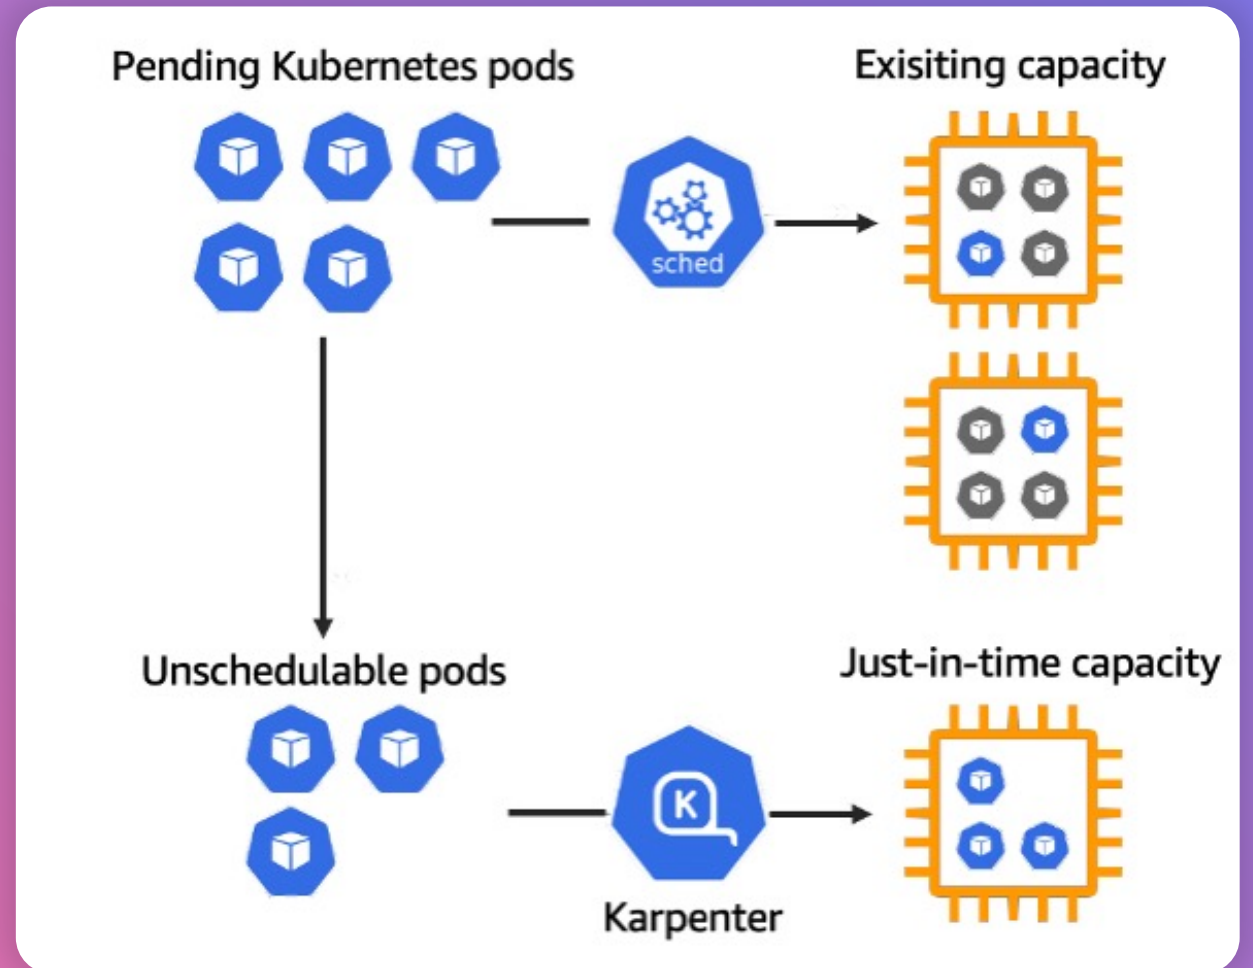
Launch template support

Custom configuration and custom AMIs for your Kubernetes nodes

Karpenter scale-up


NODE PROVISIONING

- Kube Scheduler gets the first crack at scheduling pending pods. Tries to schedule on existing capacity
- Karpenter observes aggregate resource requests of **unschedulable pods** (set by kube scheduler) to make decisions on what instances to launch




Karpenter bin-packing

NODE PROVISIONING



Online bin-packing while scaling up



Well-known labels

`karpenter.sh/capacity-type=spot`

`kubernetes.io/arch=arm64`

`topology.Kubernetes.io/zone=us-west-2a`

`node.kubererenetes.io/instance-type=m5.large`

Karpenter scale-up (continued)

NODE PROVISIONING

Instance types

Defaults to all instance types excluding metal and GPU

Diversify across sizes, families, generations CPUs

Purchase options

Defaults to on-demand

Combine Spot and on-demand

When included, Spot is prioritized

```
spec:
  requirements:
    - key: karpenter.sh/instance-type
      operator: NotIn
      values: ["m5.large"]
```

```
spec:
  requirements:
    - key: karpenter.sh/capacity-type
      operator: In
      values: ["spot", "on-demand"]
```

Karpenter scale-up (continued)

NODE PROVISIONING

Architecture

Defaults to x86 instances (amd64)

Diversify across x86 and ARM

Availability zones

Defaults to all AZs

```
spec:
  requirements:
    - key: karpenter.sh/arch
      operator: In
      values: ["arm64", "amd64"]
```

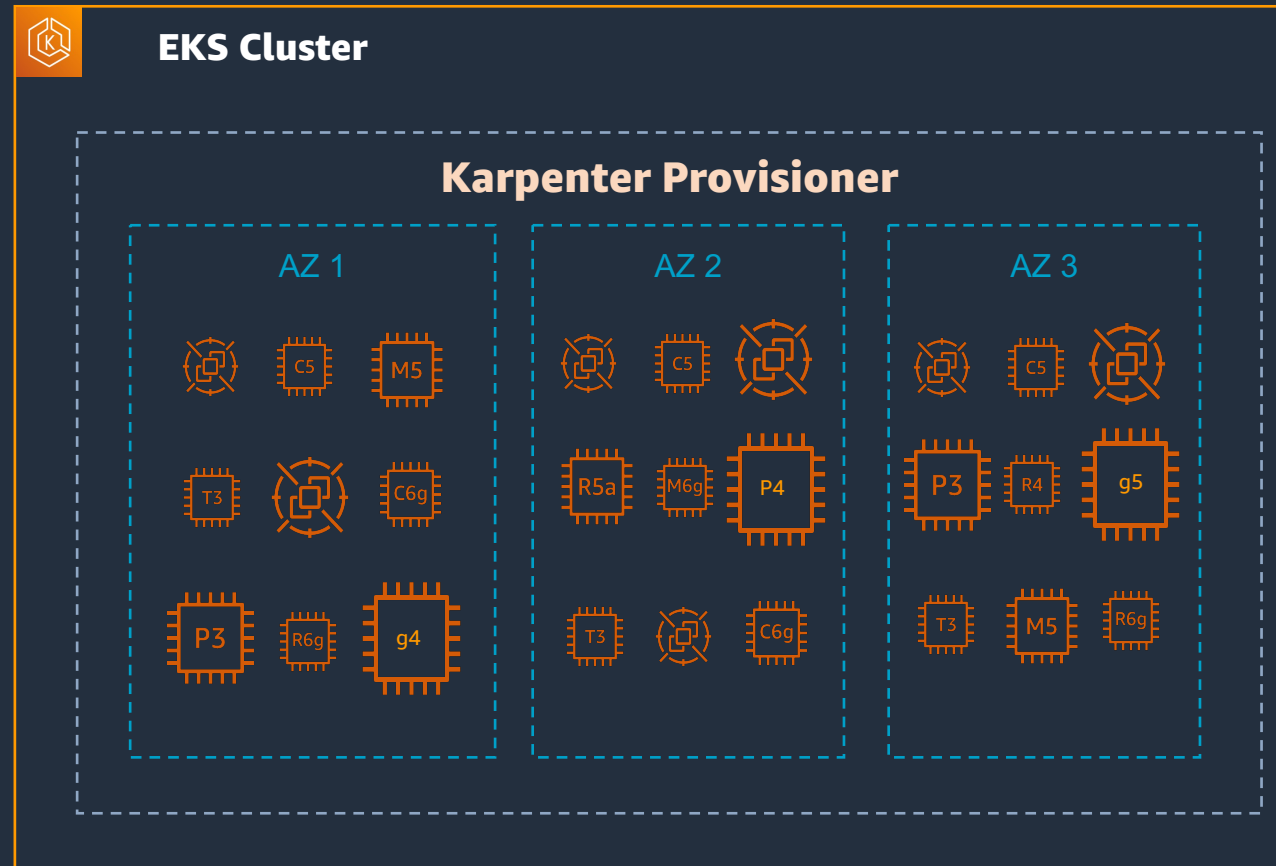
```
spec:
  requirements:
    - key: karpenter.sh/zone
      operator: In
      values: ["us-west-2a"]
```

Provisioner CRD

- Custom resource to provision nodes with a set of attributes (taints, labels, requirements, TTL)
- Single provisioner can manage compute for multiple teams and workloads
- Can also have multiple provisioners for isolating compute for different needs

```
apiVersion: karpenter.sh/v1alpha5
kind: Provisioner
metadata:
  name: default
spec:
  consolidation:
    enabled: true
  requirements:
    # Include general purpose instance families
    - key: karpenter.k8s.aws/instance-family
      operator: In
      values: [c5, m5, r5]
    # Exclude small instance sizes
    - key: karpenter.k8s.aws/instance-size
      operator: NotIn
      values: [nano, micro, small, large]
    - key: karpenter.sh/capacity-type
      operator: In
      values: ["on-demand", "spot"]
    - key: kubernetes.io/arch
      operator: In
      values: ["amd64", "arm64"]
  providerRef:
    name: default
```

Going large scale with Karpenter and Flexible Compute



Karpenter scale-in

NODE TERMINATION

Node TTL

Terminate empty nodes

Expire nodes to relaunch with new AMIs

Consolidation

Attempts to reduce the overall cost of the nodes launched by that provisioner if workloads have changed

```
kind: Provisioner
metadata:
  name: default
spec:
  ttlSecondsAfterEmpty: 30
  ttlSecondsUntilExpired: 2592000
```

```
kind: Provisioner
metadata:
  name: default
spec:
  consolidation:
    enabled: true
```

Karpenter consolidation

INTELLIGENT WORKLOAD'S RESOURCE CONSUMPTION REDUCTION


Deletes a node when...

It is empty or pods can run on free capacity of other nodes in the cluster

Replaces a node when...

Pods can run on a combination of free capacity of other nodes in the cluster + more efficient replacement node

Reduced Complexity, Increased Performance at scale



<https://www.vladionescu.me/posts/scaling-containers-on-aws-in-2022/>

Takeaways

- Schedule pods to EC2 Spot Instances to **optimize cost**
- Use **Provisioners** to ensure you are scaling nodes using Spot best practices
- Use **default Provisioner** with diverse Instance Types and Availability Zones
- Use **additional Provisioners** for different compute constraints
- Control scheduling of your application Pods with **Node Selector, topologySpreadConstraints, Taints, Tolerations** and **Provisioners**

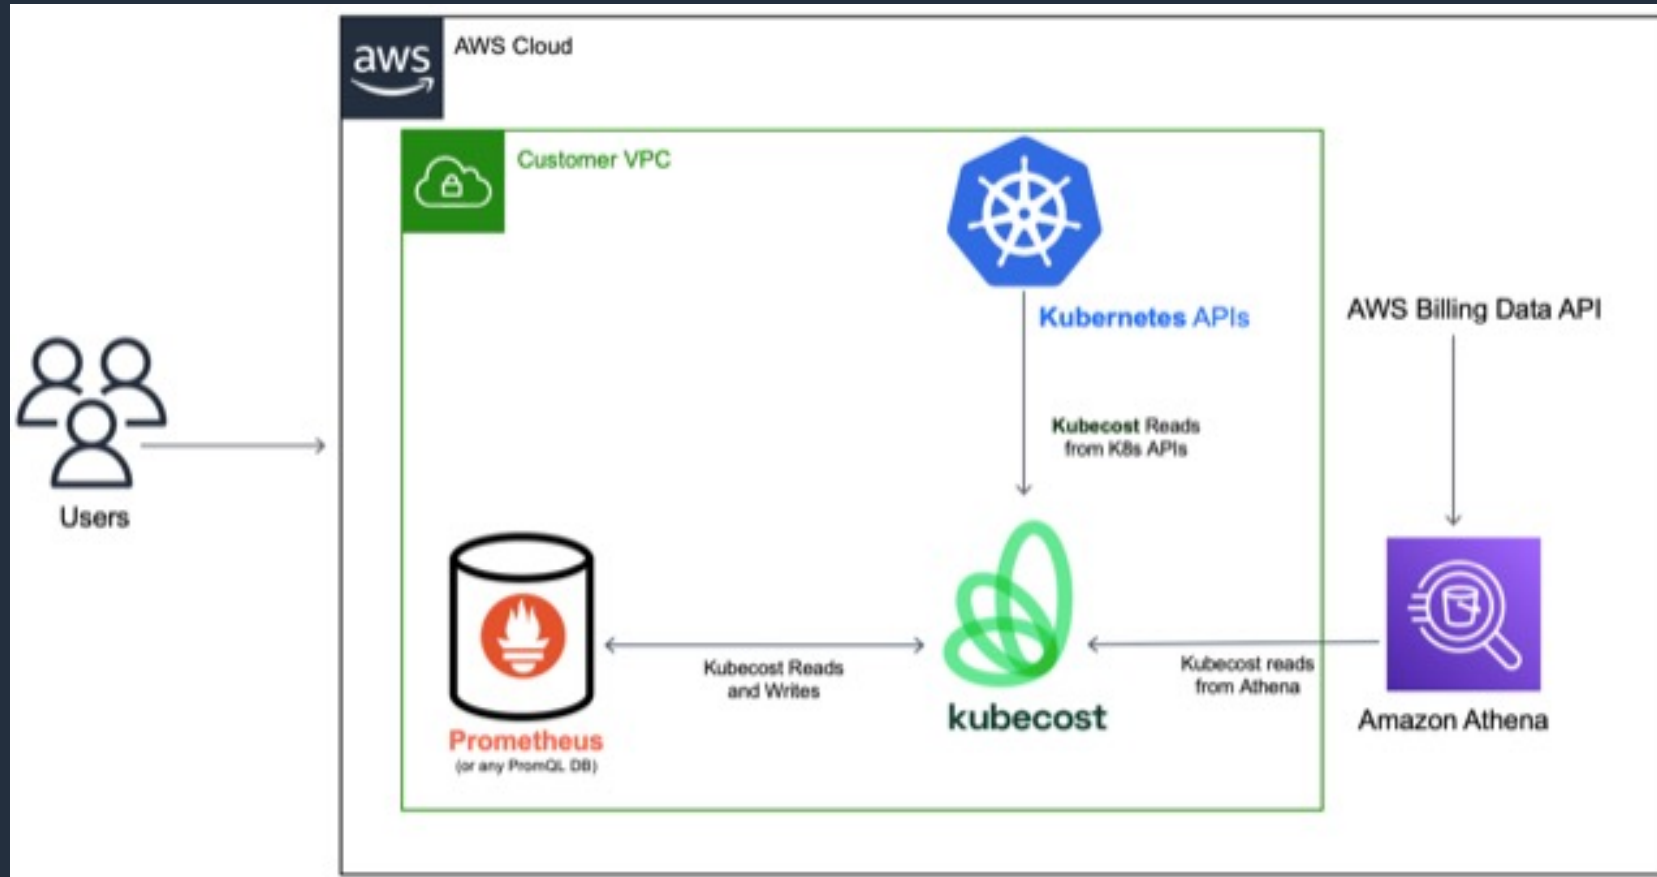
Kubecost



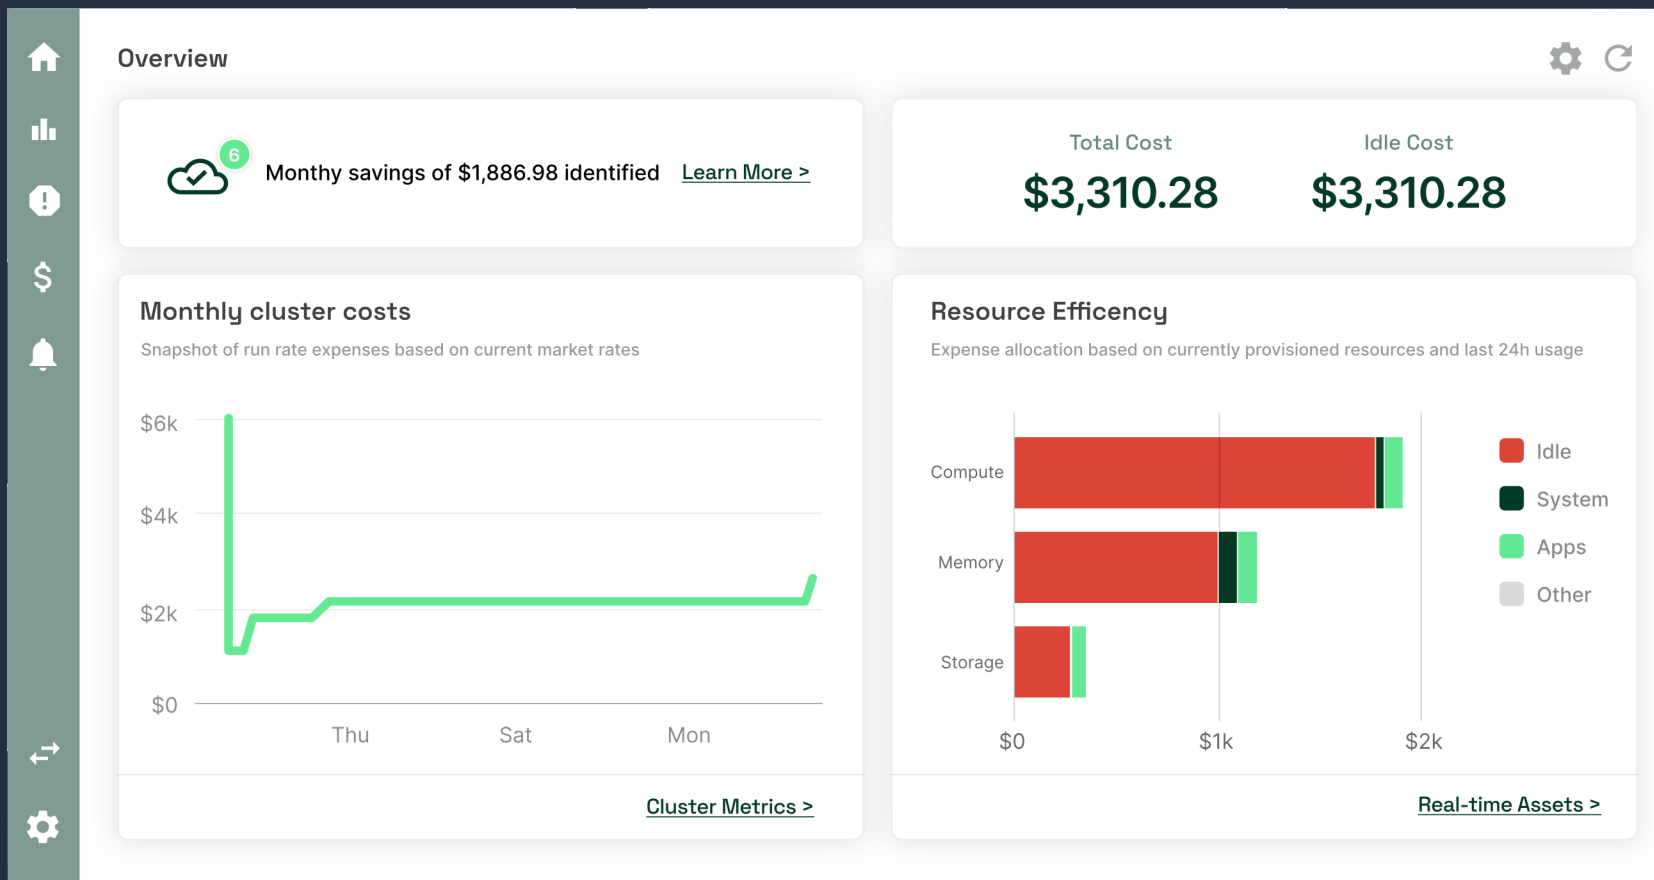
Kubecost: Real-time Cost Monitoring

- Open source tools like **Kubecost** enable real-time cost visibility for Kubernetes
- Evaluates cost and usage at deep granularity:
 - by Kubernetes service, deployment, namespace, label, statefulset, daemonset, pod, and container
- Cost and usage can be attributed to org concepts such as team or application
- Uses Prometheus metrics to determine usage by applications
- Makes recommendations for where to optimize resources

Kubecost: Real-time Cost Monitoring



Kubecost: Real-time Cost Monitoring





Thank you!

Borja Pérez

bpguasch@amazon.es

Arpit Sapra

sapraa@amazon.com

SCAN ME



Additional resources

- Karpenter : <https://karpenter.sh>
- Karpenter Best Practices: <https://aws.github.io/aws-eks-best-practices/karpenter/>
- Karpenter workshop: <https://ec2spotworkshops.com/karpenter.html>
- Launch Blog: <https://aws.amazon.com/blogs/aws/introducing-karpenter-an-open-source-high-performance-kubernetes-cluster-autoscaler/>
- Blog post: <https://aws.amazon.com/blogs/containers/using-amazon-ec2-spot-instances-with-karpenter/>